

Meridian

A Mathematical Language with Certified Computation

From readable proof steps to a domain-aware,
proof-producing computer algebra workspace

Central proposal. Write the mathematical choices; let verified algorithms construct their consequences; retain the exact domains and specifications of every computed object. A calculation is not merely an expression returned by a solver: it is an object accompanied by a precisely stated theorem about that object.

Prepared for Vladimir Reshetnikov's Brainstorm campaign

Prepared by ChatGPT

Date September 5, 2026, Pacific time

Status. This is a research and engineering design, not an implemented language. It is based on a close reading of both requested synthesis reports, static inspection of selected committed ProveIt and Leant sources, and primary literature on verified computation. The mathematical examples include explicit proofs. The accompanying Python checks were executed; they are not Lean proofs. No new Lean compiler, CAS integration, or human-subject evaluation is claimed.

Abstract

The first Brainstorm iteration reached a useful consensus: retain Lean’s checking foundation, make mathematical steps visible, elaborate omissions into scoped obligations, preserve the intended statement, and retain evidence rather than repeat open-ended discovery. A second iteration should not merely rename this architecture. This article proposes a specific extension: a mathematical language whose central objects include *certified computations*, not only proof goals. The resulting system, provisionally called Meridian, combines a structured proof surface with a computer algebra workspace built around formally specified algorithms, executable representations, and kernel-checkable evidence.

Its distinctive features are contracts that distinguish an identity from a complete solution, domain-preserving simplification, explicit branch coverage, composable representation packages, and a shared interface for verified algorithms, certificate-producing external CAS tools, Lean tactics, and Leant’s term synthesis. The main worked example derives the polynomial recurrence for all derivatives of both real Lambert branches, separating a reusable symbolic calculation from local analytic hypotheses. Further examples cover rational cancellation, complete root solving, finite binomial transforms in additive commutative groups, radical denesting, and finite jets of formal series.

The proposal resolves concrete questions raised by the two syntheses, specifies trust and replay boundaries, gives a bounded implementation plan, and supplies negative tests that distinguish mathematical validity from plausible computation. Its central empirical hypothesis is that attaching precise, reusable meaning to computations reduces both authoring friction and maintenance cost more effectively than hiding additional tactic lines. That hypothesis remains to be tested against ordinary Lean supplied with exactly the same libraries and algorithms.

Contents

1	The proposal in one argument	2
2	What the two syntheses establish—and what they leave open	3
3	A mathematical workspace, not a larger tactic vocabulary	4
4	Computation contracts: say exactly what was obtained	7
5	What a formally verified CAS would actually require	9
6	Domain-preserving computation	11
7	Definitions and representations as reusable mathematical packages	14
8	Main case study: symbolic computation inside an analytic proof	16
9	Finite algebra without accidental loss of generality	19
10	Formal series, radicals, and numerical enclosures	20
11	Leant as a partner in the hybrid system	22
12	A realizable architecture and evidence format	24
13	A bounded implementation plan	27
14	Evaluation: isolate the language from the mathematics	28
15	Decisions for the second brainstorm iteration	30
16	The remaining research questions	31
17	Conclusion	32
A	Minimal protocol and semantic interfaces	33
B	Source register and inspection boundaries	34
C	Companion calculations and limits of validation	34
	References	35

1 The proposal in one argument

A mathematician often writes a proof in a mixture of logical and computational moves: choose a substitution, rewrite a sum, differentiate an identity, solve a polynomial equation, inspect a few coefficients, conjecture a recurrence, and then prove that recurrence for every index. A proof assistant can represent all these activities, but the user interface need not force them all into the same shape of “a goal followed by tactics.” Conversely, a computer algebra system can manipulate expressions without making the precise scope of each answer conveniently available to subsequent reasoning.

I propose a workspace in which a *mathematical claim* and a *computed mathematical object with a specification* are equally natural units. The language should support both

establish P and construct y satisfying $\text{Spec}(x, y)$.

The first asks for evidence of a fixed proposition. The second asks for data and for evidence of a property of that exact data. The second operation is not safely replaced by finding an inhabitant of the output type.

For example, a returned list of roots may contain only genuine roots while omitting half the solution set. A factorization may multiply back to the input without being irreducible. An antiderivative may be correct on one component of a domain without specifying every antiderivative on the whole domain. A truncated series may satisfy an equation modulo a power of its variable without solving the equation as an infinite series. These are not peripheral exceptions. They determine what subsequent mathematics may legitimately conclude.

Design commitment. Every computational verb has an exact result contract. The language never upgrades “a verified candidate” to “the complete answer” unless the stronger property has its own evidence. Domains, branches, coefficient structures, and approximation relations belong to the contract, not to a warning printed after the answer.

The proposal retains the first-round architecture rather than replacing it. Meridian is initially a Lean-hosted mathematical interface, with an ordinary Lean frontend and a structured document frontend over shared methods. Its foundation is not a new logic. Its new work lies in the organization of computations and their specifications, in predictable transport between representations, and in the authoring experience around those operations.

This also changes what should be automated. A system should be particularly good at taking a *chosen* substitution, recurrence ansatz, or representation and producing all the consequences that follow mechanically. It may suggest a better choice, but adopting that suggestion is a visible change to the mathematical plan. A proof which is shorter because the system silently chose a different problem is not a success.

1.1 What “human-scale” should mean

The objective is not to mimic every sentence of informal prose. It is to let the reader see the invariant, the useful representation, the decisive reduction, and the reason the calculation applies. The details may then be expanded where they matter. A domain condition such as $w \neq -1$ is visually prominent in a Lambert calculation because it controls the pole; the proof that a positive exponential is nonzero can normally remain collapsed.

Exploration and exposition should differ. During exploration, the workspace can hold several candidate normal forms and competing proof plans. In the published argument, it retains the accepted object, its specification, and a concise reason for its use. The discarded candidates may remain in an optional private search log. They are not automatically part of the proof document.

Nor should the system insist on prose when a formula or a short Lean expression is already better. Existing theorem applications and calculation blocks remain first-class input. A new surface earns its place by improving the expression or understanding of a mathematical decision, not by translating every line into English.

2 What the two syntheses establish—and what they leave open

2.1 A close-reading position

The two requested reports are complementary. *Mathematical Ideas, Checked Evidence*, here called synthesis A, develops a six-commitment common architecture, careful acceptance distinctions, a substantial static review of Leant, and specific next-round experiments. *Nine Blueprints for a Mathematician’s Proof Language over Lean*, synthesis B, supplies a finer feature matrix, a lexical corpus audit, the annotated-theorem proposal, and a strong challenge to the idea that new syntax is the main missing ingredient [1, 2].

The most important shared conclusions should be treated as the starting point, not as this article’s innovations: a fixed resolved statement; ordered dependent contexts; explicit mathematical choices; reusable theorem-backed methods; independent checking of candidates; evidence attached to all asserted claims; and fair comparison with refactored Lean. The individual proposals are discussed here through the syntheses’ detailed crosswalks, rather than presented as nine newly replicated implementation studies [1, 2].

There are also places where a second-round design should be more precise than an attractive slogan. A statement seal is not necessarily obtained by one isolated elaboration pass: types, instances, and universe constraints can remain interdependent. The meaningful requirement is a checked, immutable semantic boundary before an external proof producer receives authority to solve the request. An abstraction-safety flag is useful evidence *about a search procedure*, but is not itself a kernel proof of a negative mathematical claim. A measured frequency of cast-like lines is not a measurement of removable work. These distinctions govern the design below.

2.2 How this proposal differs from another first-round blueprint

Starting point in the syntheses	Additional decision in Meridian
Typed assertions and an obligation ledger	Computed objects also carry contracts: identity, domain equivalence, completeness, canonicity, or a specified error relation. These contracts are not collapsed into a single status.
Certified representation views	A compositional interface fixes direction, guards, semantic observations, and whether a route preserves or reflects a property. Competing meaning-bearing routes require an explicit choice.
Definition authoring as an open problem	Start with definition packages for partial expressions and executable polynomial representations, not a general natural-language definition parser.
Methods as annotated theorems	Use ordinary theorems for logical assembly; use verified algorithms or verified certificate checkers for computational families. Give both the same result-contract interface.

Starting point in the syntheses	Additional decision in Meridian
Leant behind a typed boundary	Separate candidate generation, exact displayed-edit replay, mathematical specification proof, and whole-document completion.
Retained evidence	Archive the exact target, accepted data, certificate inputs, checker theorem, proof evidence, and environment. A computed answer is part of the mathematical artifact, not just a tactic trace.
Shared-method evaluation	Add an algorithm-versus-interface ablation: a better CAS backend and a better language must not receive credit for the same gain.

2.3 What the repository examples actually show

Three directly inspected ProveIt modules are especially informative. `BinomialInversion.lean` separates integer-kernel algebra from transforms on an arbitrary additive commutative group. Its final results already reuse a triangular-kernel framework. `AutonomousIteratedDeriv.lean` isolates a local chain-rule theorem whose differentiability assumptions apply only at values actually reached by the solution. `LambertWHigherDerivatives.lean` combines that theorem with explicit exponential, polynomial, and quotient-rule calculations [4, 3, 5].

The third module supplies the central new test in this article. The symbolic calculation and the local analysis are already mathematically separated, but there is an opportunity to give the former a reusable executable interface. The new language should not hide or weaken the latter. This is a stronger test than making an already available final theorem easier to invoke.

Synthesis B’s audit classifies roughly one third of its tactic-like source entries into families proposed for reconstruction, and synthesis A explicitly warns against interpreting that figure as a savings estimate. I adopt that warning. This article does not rerun the audit, claim that the sampled corpus was compiled, or infer a percentage reduction in development time from source frequencies [1, 2].

2.4 Source identity and scope of this review

Both syntheses were read at Brainstorm revision 494be9d5. The public refs resolved during this review were ProveIt 24ce8bd7 and Leant 3a40904b. The syntheses also discuss a different local ProveIt snapshot and uncommitted Leant work. Those observations remain attributed to the reports; they are not silently recast as observations of the committed files inspected here. Full revision identities and file locators appear in Appendix B.

3 A mathematical workspace, not a larger tactic vocabulary

3.1 The four kinds of persistent node

The proposed intermediate representation has four principal kinds of mathematical node. A *claim* has a fixed proposition and evidence. An *object* has a fixed type, a chosen value, and its declared specification. A *computation* links input objects to output objects through a checked relation. A *scope* binds parameters, assumptions, definitions, notation, and local representation policies. An operation may create several nodes at once.

The dependency structure is more than a graph of propositions. For example, choosing a polynomial order affects a computed normal form; choosing a root affects later functions;

choosing a domain changes which substitutions are legal. Those are data dependencies. They must be recorded even when two choices happen to make the same proposition provable.

This gives a single workspace two useful directions. A backward request such as “show this expression is nonnegative” asks for a certificate meeting a fixed goal. A forward request such as “factor this polynomial over $\mathbb{Q}$ ” produces a candidate factor object and a theorem about it. A later proof can use that object without rediscovering it or reparsing its printed form.

3.2 An inspectable semantic header

Each declaration has a resolved header recording its binders, dependency order, constants, instances, and chosen interpretations. The reader-facing view should show the few choices that distinguish nearby mathematical statements: natural versus integer subtraction, total versus partial division, formal versus analytic series, and pointwise versus neighborhood equality.

This is not an algorithm for recovering the author’s intention. A deterministic rendering can disclose what was formalized; the author must still recognize whether it is the desired statement. The renderer and source map need their own tests. A hash of an elaborated expression identifies a record but does not make the record intelligible.

Local conventions are lexical and explicit. For example, an author can declare that an expression block uses partial real functions while the surrounding Lean file continues to use its ordinary total operations. The system then inserts proved interface lemmas at the boundary. It must not change the meaning of existing Lean notation globally.

3.3 A small proposed surface

The following is *proposed syntax*, not executable Lean or an implemented Meridian parser:

```
context rational_functions:
  variable x : Real
  let f := partial (x^2 - 1) / (x - 1)

  calculate g := simplify f
    preserving domain
    by rational_normal_form

  have roots : solutions(f = 0) = {-1}
    by solve_completely over Real

  explain "Cancellation preserves the original missing point."
```

The first command creates a mathematical object with a specified domain. The calculation binds a new representation of the same partial function. The claim about solutions requires a complete solution-set contract. The final line is ordinary commentary, visibly distinguished from a formally asserted claim.

For a logical argument the surface stays close to a familiar structured proof:

```

proof:
  fix n
  induct n maintaining the identity on U
  case zero:
    conclude by the defining equation
  case successor n:
    fix x in U
    differentiate the induction identity near x
    conclude by the chain rule and the recurrence

```

The maintained identity, the open set U , and the recurrence are bound typed objects, not strings that a model resolves by plausibility. The short English aliases are a presentation layer over stable references.

A narrow first grammar needs declarations, named claims, object bindings, calculations, scopes, case splits, induction with a stated motive, and a Lean escape block. It does not need unrestricted English theorem statements. The extension beyond the syntheses is primarily the semantics of `calculate` and of its output, not a proliferation of keywords.

3.4 Author-selected inputs versus reconstructible requirements

The logical contract of a method remains an ordinary theorem where possible. Its binders receive roles such as *input*, *choice*, and *requirement*. The distinction is contextual. A bijection can be routine in one reindexing and the central idea in another. A “requirement” is never an assertion that its proof is easy.

For long-term stability, I would register roles on a small project-owned wrapper with a documented semantic interface. That wrapper may delegate to a changing Mathlib theorem. This creates a modest, explicit maintenance cost, but avoids making every upstream binder rename a language-API change. The wrapper’s checked type is authoritative; operational metadata is separately versioned. Positional binder lists and automatic classification of every proposition as a side condition are both too fragile.

Named methods constrain the accepted construction when their wording asserts a method: “by reindexing” must actually supply a reindexing contract. An explicitly open justification, such as “by search,” permits another proof of the same claim. An alternative argument is welcome, but accepting it changes the displayed reason rather than leaving an inaccurate explanation in place.

3.5 How the interface supports mathematical thought

Suppose an author guesses a form $e^{-nw}P(w)/(1+w)^k$. The workspace can immediately show its derivative, the denominator conditions, and the polynomial factor needed to close the induction. The author sees whether the ansatz is stable under the differential operator. This is a much more useful interaction than receiving only a sequence of suggested quotient-rule tactics.

Similarly, a sum can be displayed as a function of its index region, with a chosen map and its image. A failed reindexing can show the missing endpoint. A formal series calculation can display exactly which coefficients have been determined. The proof state remains available, but it is no longer the only representation of progress.

Analogy is treated as checked reuse. “Similarly for the lower branch” instantiates a known template with different branch hypotheses; the newly required inequality is visible. Induction is treated as a mathematical invariant rather than an instruction to manipulate whatever goal happened to be active. These mechanisms do not certify a human mental process. They preserve a useful mathematical explanation of the accepted argument.

4 Computation contracts: say exactly what was obtained

4.1 A result is a value together with a proposition about it

For input $x : I$, output type $O(x)$, and a specification $\text{Spec}(x, y)$, the ideal mathematical result has the dependent type

$$\text{Result}(x) = \{y : O(x) \mid \text{Spec}(x, y)\}.$$

An implementation may store a certificate from which the proof is reconstructed rather than a directly serialized dependent pair. Nevertheless, the acceptance boundary must establish the displayed specification for the *exact* accepted output. Verifying a similar expression, an earlier candidate, or the output of a different renderer does not establish it.

There is no single universal specification for “simplify,” “solve,” or “factor.” The interface needs a small family of explicit contracts. A result can carry several compatible properties, each with its own evidence. The properties form a partial order under proved implications, not a universal ladder on which every task receives one numerical confidence score.

Request	Minimal useful certificate	Additional property not implied
Rewrite an expression	Equality under a stated domain and interpretation	Equality of domains; preferred form; minimal expression size.
Factor a polynomial	Product of returned factors equals the input	Irreducibility, uniqueness conventions, or completeness over another field.
Find solutions	Every returned object satisfies the equation	Every solution is returned; multiplicity; a canonical ordering.
Solve completely	Membership in the returned set is equivalent to the original predicate	Efficient membership, finite enumeration, or preferred parametrization.
Compute a derivative	A derivative assertion on the stated domain	Maximal differentiability domain or validity at excluded boundary points.
Integrate symbolically	The returned function has the requested derivative on a domain	A definite integral formula or all antiderivatives on disconnected domains.
Compute a series jet	Equality of coefficients below a fixed order	An infinite identity, convergence, or an analytic remainder estimate.
Enclose a value	Exact lower and upper bounds	Exact equality to a decimal or minimal interval width.
Prove ideal membership	An exact identity $p = \sum_i h_i f_i$ in the specified coefficient ring	A Gröbner basis, non-membership, or ideal membership over a smaller coefficient ring.
Select a root	Its defining equation and selected branch condition	Uniqueness without an additional theorem; executable selection from existence alone.

These distinctions are mathematical rather than cosmetic. Davenport’s study of verified polynomial factorization explicitly separates checking a product identity from verifying irreducibility [15]. Meridian should make that separation unavoidable in the result type, not merely repeat it in documentation.

4.2 Complete solutions are predicates before they are lists

Let $P : A \rightarrow \text{Prop}$ be the original solution predicate. A returned list L can be certified by

$$\forall y, y \in L \implies P(y).$$

A complete finite answer needs the stronger property

$$\forall y, y \in L \iff P(y).$$

A system solving $ax = 0$ over $\mathbb{R}$ cannot always return a finite list: when $a = 0$ every real number is a solution. The result type should therefore permit a symbolic set or a parameterized family, with a theorem characterizing its membership. Enumeration is an additional representation choice.

For the parameterized example, a complete answer is

$$\{x \in \mathbb{R} : ax = 0\} = \begin{cases} \mathbb{R}, & a = 0, \\ \{0\}, & a \neq 0. \end{cases}$$

In the second case division by a proves necessity, and substitution proves sufficiency. In the first case multiplication by zero proves every x is a solution. The cases cover all a . A generic answer “ $x = 0$ ” with the unproved condition $a \neq 0$ is useful, but it is a *conditional answer*, not the complete parameterized solution.

Multiplicity is a separate contract. A set of polynomial zeros intentionally forgets multiplicity; a multiset factorization does not. Over $\mathbb{R}$ and over $\mathbb{C}$ the zero sets may also differ. The solver does not infer a field merely because one field produces a simpler answer.

4.3 Canonicity must be defined, not suggested by typography

An algorithm can return any correct factor order unless an order is required. An algebraic number can have many defining polynomials and isolating intervals. A normal form is canonical only relative to specified conventions: coefficient normalization, variable order, term order, algebraic extension, and an equality notion. The interface should never display “canonical” simply because the output is deterministic in one run.

For most proof steps, canonicity is unnecessary. A proof needs a convenient representative with a transport theorem. Demanding a global canonical form can increase checking cost and make the interface less stable. Meridian therefore asks for the weakest sufficient mathematical result, while retaining the opportunity to request stronger properties explicitly.

4.4 Finite observations versus universal behavior

A synthesized function may satisfy a test at a few inputs while failing its intended universal law. Leant’s separation between term construction and behavioral assessment is directly relevant here. Its committed verification wrapper explicitly records callback acceptance rather than a kernel or behavioral certificate [6]. The syntheses also describe an uncommitted candidate-bound proposition check; that observation is attributed to those snapshots, not asserted as an API executed in this review [1, 2].

Meridian records the proposition actually established. A proof of a finite test conjunction certifies that conjunction. A proof of a universal law certifies the law. Some universally quantified propositions over finite types are decidable; the distinction is not the presence of a quantifier but the exact specification and the evidence returned. A provisional data object can be explored under a weaker tested property, but a published theorem cannot quietly consume the unproved stronger one.

5 What a formally verified CAS would actually require

5.1 Three implementation lanes, one logical interface

The hybrid system should support three lanes while labeling their guarantees honestly.

Verified algorithm lane. A function $a : I \rightarrow O$ is defined in the host logic, and a theorem proves

$$\forall x, \text{Spec}(x, a(x)).$$

This is the appropriate target for stable computational building blocks: polynomial arithmetic, a syntax-directed derivative algorithm on a specified expression grammar, exact rational arithmetic, and truncated series operations. Termination is established for the host definition or separately for the stated algorithmic interface.

Certified external-result lane. An untrusted program returns (y, c) . A host-defined checker check and its soundness theorem establish

$$\text{check}(x, y, c) = \text{true} \implies \text{Spec}(x, y).$$

The search or factorization algorithm need not be verified. The accepted result is nevertheless rigorous when the checker theorem and its actual application are checked. This is a verified *checking boundary*, not a claim that the external CAS algorithm has been formally verified.

Proof-producing symbolic lane. A tactic or compiler constructs a proof of $\text{Spec}(x, y)$ directly using registered lemmas. A derivative expression can be built recursively together with derivative evidence. Algebraic cleanup can then use an existing proof-producing normalizer. Such a producer may remain logically untrusted because the resulting term is checked at the exact requested type. Its operational execution still requires appropriate isolation.

All three lanes can inhabit the same result contract. They differ in what is proved about the program and in how a concrete execution acquires authority. A practical system should use verified algorithms for its reusable core and allow certificate-producing external services for expensive search. It should not delay all useful interaction until an entire general-purpose CAS has been reimplemented and verified.

5.2 The execution gap in “the algorithm was proved correct”

Suppose a verified function a is compiled and an executable prints the value y . The theorem $\text{Spec}(x, a(x))$ alone does not prove $\text{Spec}(x, y)$. One also needs an accepted connection between the concrete returned data and $a(x)$. There are several valid routes: kernel reduction identifies the values; a proof-producing evaluator constructs the equality $a(x) = y$; or an independent checker establishes the required specification directly from x, y, c .

This small distinction is central to a genuinely rigorous hybrid. Proof annotations erased by a compiler do not make arbitrary executable output self-authenticating. A formally verified source algorithm and a trusted native execution mode are different guarantees. A deployment may intentionally trust its compiler, but it must not describe that choice as the same trust boundary as kernel-only checking.

Lean’s current documentation makes the distinction observable: native computation can introduce additional axiom dependencies, and current `native_decide` uses a per-invocation axiom. A blacklist of one historical axiom name is therefore not an adequate publication policy [10]. The strict Meridian profile requires an explicit transitive allowlist. A separate performance-oriented profile may permit additional computation assumptions, but the exported artifact and reader view must state that choice.

Current Lean documentation also describes `cbv` as a proof-layer, call-by-value-style simplification mechanism. This provides a relevant candidate for evaluation of some verified computations

without equating native execution with proof. Availability and performance must be tested at the pinned project toolchain; the moving manual is not a compatibility guarantee [11].

5.3 Reification is part of the proof boundary

Efficient computation usually uses a representation different from the library’s semantic expression. Let Rep be a computable polynomial representation and $\text{denote} : \text{Rep} \rightarrow \text{Polynomial}(R)$ its interpretation. A certificate about the representation becomes useful only after the bridge to the original input has been established:

$$\text{denote}(r) = p, \quad \text{check}(r, s, c) = \text{true}, \quad \text{check_sound} : \dots \implies \text{Spec}(\text{denote}(r), \text{denote}(s)).$$

The importer must not trust that a printed variable name, coefficient string, or CAS symbol happens to denote the intended Lean object. Variables need typed identities; rational coefficients need exact numerators and denominators; noncommutative expressions cannot be flattened into commutative polynomials. A reification error is harmless only if the bridge to the fixed input is also checked and the final result is checked against the original specification.

Useful sharing does not require an enormous universal symbolic-expression type. Small domain-specific representations are preferable. A polynomial syntax can have a complete, simple semantics. A derivative syntax can include a limited set of differentiable operations and treat other functions as opaque atoms with explicit derivative hypotheses. Failure to recognize an atom produces a located unsupported result, not an invented identity.

5.4 What existing work can be reused

There is substantial prior art for the computational half. Isabelle’s Archive of Formal Proofs contains verified univariate polynomial factorization, including the Berlekamp–Zassenhaus development and associated square-free and rational interfaces [17]. Rocq’s Interval package supplies automated real inequality and enclosure reasoning, including supported integral forms [18]. These establish useful precedents, not drop-in Lean certificates: crossing foundations requires a proved translation, a Lean reimplement, or Lean-side checking of exported certificates.

For Lean specifically, the April 2026 preprint by Shen, Guo, Liu, and Zhi describes a computable polynomial interface and certificate-based integration with external CAS tools for ideal and Gröbner-basis reasoning [16]. It should be investigated as reusable work, not described as a feature of Meridian already implemented. The older `polyrith` workflow is also instructive, but current Mathlib documentation marks its external service as defunct and the tactic as unsupported [14]. A second-round implementation plan must verify an actual backend rather than rely on an old list of tactic names.

The proposed contribution is consequently not the invention of reflective computation or declarative proof documents. Isar already exemplifies structured proof documents and separates proper document constructors from exploratory commands [19]. Meridian’s contribution is a particular integration contract for this campaign: reusable computed objects, their domains and result strength, and their place in a readable Lean-checked argument.

5.5 A small verified kernel of computation, not a second logical kernel

The first computational package should cover exact scalar arithmetic, sparse polynomial operations, domain-aware rational normalization, and a small syntax-directed derivative compiler. Calling this a computational core must not confuse it with a trusted logical kernel. Its correctness theorems are ordinary Lean theorems; bugs in its untrusted callers are contained by exact checking.

The core should also expose *resource behavior*. Termination of a normalizer does not imply interactive speed. Output can expand dramatically; a compact input can have a large normal form.

Requests need size limits as well as time and memory limits. A computation that stops before completing must not issue the same result variant as a computation that proved completeness.

6 Domain-preserving computation

6.1 Three meanings of the same printed quotient

Consider

$$f(x) = \frac{x^2 - 1}{x - 1}.$$

As an element of the rational-function field $\mathbb{R}(X)$, it equals $X + 1$. As a partial real expression, it is defined only when $x \neq 1$, and agrees with $x + 1$ on that domain. As a total real expression with division by zero defined to give zero, it takes the value 0 at $x = 1$, whereas $x + 1$ takes the value 2. These are different mathematical objects.

The language must therefore distinguish these three tasks:

```
calculate in RationalFunction(Real):
  (X^2 - 1) / (X - 1) = X + 1

calculate as PartialFunction(Real, Real):
  simplify (x^2 - 1) / (x - 1) preserving domain

have for x : Real with x != 1:
  (x^2 - 1) / (x - 1) = x + 1
```

Only the first is a field-of-fractions identity without a pointwise domain. The second returns a restricted representative. The third proves a conditional equality of ordinary real values. An interface that prints all three simply as “ $x + 1$ ” has erased information that later steps may need.

6.2 A small semantic model

For fixed types X, A , model a partial expression as

$$E = (D_E, v_E), \quad D_E : X \rightarrow \mathbf{Prop}, \quad v_E : \prod_{x:X} D_E(x) \rightarrow A.$$

This is a mathematical specification, not a demand that every expression in a Lean file be replaced by a partial-function type. It is sufficient to implement it inside an explicitly selected expression package.

Definition 6.1 (Exact equivalence and guarded equivalence). *Two partial expressions E, F are exactly equivalent, written $E \simeq F$, when*

$$\forall x, D_E(x) \leftrightarrow D_F(x)$$

and their values agree for all corresponding domain witnesses. For a guard $G : X \rightarrow \mathbf{Prop}$, write $E \simeq_G F$ when the same conditions hold at all x satisfying $G(x)$.

A weaker relation, equality of values on a chosen set, does not assert equality of domains. Meridian must retain the name of the relation in the certificate. It is a type error in the computational protocol to consume an `EqOn` certificate where exact partial-function equivalence is required.

The simplified version of the quotient above should therefore be

$$F = (\{x : x \neq 1\}, x \mapsto x + 1).$$

The polynomial $x + 1$ with domain all of $\mathbb{R}$ is a different object. It may be introduced as an *extension*, accompanied by a theorem about agreement on the old domain. Extending a function is a legitimate mathematical operation, but it is not domain-preserving simplification.

6.3 Composition and specialization rules

Proposition 6.2 (Composition of guarded computations). *If $E \simeq_G F$ and $F \simeq_H K$, then $E \simeq_{G \wedge H} K$.*

Proof. At a point satisfying both guards, compose the two domain equivalences. A domain witness for E transports to one for F and then one for K . The two value equalities compose by transitivity. The required equality is independent of which domain witnesses are used, since each hypothesis asserts equality for corresponding witnesses. In Lean, proof irrelevance also makes proof-valued domain arguments interchangeable. $\square$

The same ambient variable is essential in that formulation. When a later method uses a transformed object or a substituted variable, its guard must be instantiated at that object or variable; simply collecting uninstantiated strings in a set is not composition.

Proposition 6.3 (Specialization preserves guards). *Let $\sigma : Y \rightarrow X$. If $E \simeq_G F$, then their pullbacks along σ satisfy*

$$\sigma^* E \simeq_{G \circ \sigma} \sigma^* F, \quad D_{\sigma^* E}(y) = D_E(\sigma(y)).$$

Proof. Apply the original domain equivalence and value equality at $x = \sigma(y)$. Their prerequisite is exactly $G(\sigma(y))$. No other hypothesis is introduced. $\square$

This elementary rule gives an important operational guarantee. Specializing a computation certified under $x \neq 1$ to $x = 1$ leaves the impossible prerequisite $1 \neq 1$; it does not produce a numerical answer. The default interface should report “outside the retained domain,” pointing to the cancellation step that introduced the guard.

6.4 Conditional results are useful, but remain conditional

A computation may discover a sufficient guard that the current context cannot prove. Its result can still be retained as

$$G \implies \text{Spec}(x, y),$$

with status *certified conditionally*. It is not an open theorem masquerading as a completed unconditional claim: the certified proposition is explicitly the implication. To use $\text{Spec}(x, y)$ in the original context, the user must establish G , enter a justified branch where G holds, or change the theorem statement visibly.

Guards need not be weakest possible conditions. For instance, a normalizer may require a product of denominators to be nonzero although a more subtle argument could avoid one cancellation. The system should label the result as conditional on a *sufficient* condition and allow a different method. It must not report that the original theorem is false merely because its chosen guard is false.

There should be no invisible command equivalent to “assume all generated side conditions.” Generated requirements are proof debts. An author can explicitly state a conditional theorem, but that action changes the public contract.

6.5 Complete branch results

Suppose a computation on a domain D returns branches $B_1, \dots, B_m$, with correct results on $D \cap B_i$. To claim a result covering the original domain, it needs

$$\forall x, D(x) \implies \bigvee_{i=1}^m B_i(x).$$

Disjointness is not inherently necessary. If branches overlap and each result represents the same original object there, their compatibility follows from those correctness theorems. For results selected for a more general relation, compatibility may require a separate proof.

A symbolic case object is not automatically an executable function. Execution also requires decidable branch tests or a supplied selection procedure. An existence proof obtained classically can be perfectly adequate for a theorem without being an algorithm for selecting a branch. The interface records which kind of object was constructed.

The simple parameterized equation $ax = 0$ is a good acceptance test: two branches, one coverage proof, and distinct result representations. Later examples may use root isolating intervals, algebraic discriminants, or connected components. The semantics should not need to change.

6.6 A fully worked complete-solving example

For the partial function $f(x) = (x^2 - 1)/(x - 1)$ with domain $x \neq 1$, define

$$P(x) \iff \exists h : D_f(x), v_f(x, h) = 0.$$

This says that x is in the domain and that its defined value is zero; no undefined value is evaluated outside the domain. Domain-preserving simplification gives $f(x) = x + 1$ at every admissible x . Consequently

$$P(x) \iff x \neq 1 \wedge x + 1 = 0 \iff x = -1.$$

For necessity, $x + 1 = 0$ implies $x = -1$. For sufficiency, $-1 \neq 1$ and $(-1)^2 - 1 = 0$, so substitution gives $f(-1) = 0$. This proves completeness, not just soundness of the returned root.

A solver which clears denominators first obtains candidate roots -1 and 1 of $x^2 - 1 = 0$, then must filter the candidates by the original domain. If instead the input were the *total* real expression with the convention above, 1 would indeed be an additional solution because the expression there is $0/0 = 0$. The same printed equation can therefore have different complete answers under different, legitimate conventions. This is why a statement view and domain-aware objects are necessary even when every individual algebraic tactic is sound.

6.7 Conditional rewriting cannot leak through a global cache

A simplifier may internally maintain an equality graph or cache normal forms. An edge established under $x \neq 1$ must remain conditional or belong to a context where that assumption has been proved. It cannot merge two global equivalence classes and then forget the guard. Reuse after undo, substitution, or a change of scope must supply a checked context map.

Meridian's initial implementation should use a simpler discipline than a global conditional equality engine: local normalizers with explicit assumptions, immutable input identities, and cached certificates parameterized by their actual context. Broader equality saturation is a later optimization. Correctness does not require solving that engineering problem first.

7 Definitions and representations as reusable mathematical packages

7.1 A definition package has a semantic center

The syntheses correctly identify definitions as a major source of downstream friction [1, 2]. Meridian should begin with a small number of *definition packages*. Each package contains a semantic type or object, one or more executable representations, denotation maps or representation relations, invariants, proved interface lemmas, and a scoped normalization policy. The package is also available to ordinary Lean without the new surface.

A polynomial package, for example, can use sparse integer coefficients for execution while exposing semantic polynomials over $\mathbb{R}$ through a coefficient homomorphism. A partial-expression package can expose a domain and a value, with operations for restriction, extension, and exact equivalence. A series package can expose finite jets with a precise projection from a formal series. These packages attack recurring representation choices without attempting to parse arbitrary new definitions from English.

An object need not carry every proof eagerly. The package can store a compact representation and reconstruct invariant evidence when needed. The invariant must still be checked before a method relies on it. Memoization changes when checking happens, not which theorem must be established.

7.2 A concrete composition interface

For representations A, B of a semantic space S , let $\text{denote}_A : A \rightarrow S$ and $\text{denote}_B : B \rightarrow S$. A view may return $b : B$ together with

$$G_1(a, b) \implies \text{denote}_A(a) = \text{denote}_B(b).$$

A subsequent view returning $c : C$ supplies

$$G_2(b, c) \implies \text{denote}_B(b) = \text{denote}_C(c).$$

The composed certificate retains the *accepted intermediate object* b and the instantiated guard

$$G_1(a, b) \wedge G_2(b, c).$$

Its semantic equality is ordinary transitivity. Where types or observations are dependent, the interface uses the corresponding dependent transport theorem; plain equality in one fixed S is the first implementation target, not a claim to cover every dependent representation automatically.

Obligations are emitted in dependency order. If the second view needs a property of b derived by the first, it receives that proof explicitly. It may not use a conclusion that will exist only after the final target is solved. A dependency cycle is an unfinished computation, not a reason to guess that all guards hold.

7.3 A three-view stress test

Take natural numbers a, b and consider their difference. The route

$$\mathbb{N} \longrightarrow \mathbb{Z} \longrightarrow \mathbb{Q} \longrightarrow \mathbb{R}$$

preserves the values of individual natural numbers. But the equality

$$\text{cast}(a - b) = \text{cast}(a) - \text{cast}(b)$$

requires $b \leq a$, where $\dot{-}$ denotes truncated natural subtraction. All later casts must retain this requirement. At $a = 2, b = 3$, the two sides would otherwise be 0 and -1 .

Now add a competing route that represents natural subtraction by $\max(a - b, 0)$ over the reals. This route can preserve the value without the order guard. It does not justify replacing the result by $a - b$. Both routes can be correct, but they expose different guards and expressions. Meridian records the chosen route; adding the new route must not silently reinterpret old source. The system may propose it as a mathematically visible simplification of the requirements.

Exact division provides a related test. Transport of a natural quotient to a rational quotient needs the relevant exactness condition, rather than the assumption that all division symbols denote field division. For a mathematical exact-division operation, a nonzero divisor and a divisibility proof form an appropriate explicit contract. At $3/2$, the natural quotient is 1, not the rational number $3/2$. When a library theorem has a weaker or differently stated premise, its actual checked type, rather than this illustrative contract, is the authority.

7.4 Preservation and reflection are different capabilities

A ring homomorphism carries a polynomial identity to an identity in the target ring. It need not be injective for that use. Conversely, proving equality after a homomorphism and concluding equality before it requires an appropriate reflection theorem, such as injectivity. Order comparisons require their own preservation or reflection facts.

This distinction matters for ProveIt’s binomial inversion: an integer identity can be transported to a commutative ring without assuming characteristic zero, and integer coefficients can act on an additive commutative group without making that group into a field [4]. A CAS backend which insists on field coefficients must not cause the public theorem to acquire an unnecessary field hypothesis.

Even an injective coefficient extension does not reflect every algebraic property. In $\mathbb{Q}[X]$, X belongs to the ideal generated by $2X$, because $X = \frac{1}{2}(2X)$. In $\mathbb{Z}[X]$ it does not: the coefficient of X in $(2X)H(X)$ is even for every integer polynomial H , and cannot equal 1. Thus a certificate of ideal membership over $\mathbb{Q}$ cannot be relabeled as an integer ideal-membership certificate. The inclusion of coefficients is injective, but the property being transported quantifies over coefficient witnesses that have changed domains.

Likewise, “nonzero” is not a universal substitute for “invertible.” In $\mathbb{Z}/6\mathbb{Z}$, the nonzero element 2 satisfies $2 \cdot 0 = 2 \cdot 3$, so cancellation fails. Contracts must name the exact algebraic structure and capability required by the transformation.

7.5 Coherence and package ownership

For the initial system, automatic view composition should use a unique explicit route selected by the local package. If more than one meaning-bearing route matches, the author selects one, or the package supplies a checked coherence rule and a documented deterministic policy. Search success is not a route selection rule.

Two packages may share notation while disagreeing about division or normal forms. Import order must not decide the mathematical meaning. Semantic conflicts require qualification or a lexical selection. Proof-search preferences may be combined under a documented budget, but choices affecting accepted data or statement interpretation remain explicit.

A definition package should have an owner, a versioned interface, migration tests, and a small collection of negative neighbors. A change to its semantic object requires a meaning-aware migration; a change only to the normalization algorithm may reuse a contract if new outputs are independently checked and accepted data are preserved or visibly replaced. This turns library maintenance into a manageable engineering boundary instead of a collection of implicit global

conventions.

8 Main case study: symbolic computation inside an analytic proof

8.1 The actual ProveIt decomposition

The inspected Lambert module defines a polynomial family, a corresponding family of real functions, a derivative family, and a scalar autonomous equation. It proves the derivative calculation and the recurrence separately, then invokes `AutonomousODE.iteratedDeriv_eq_comp` for the principal and lower real branches. The branch-specific work establishes the first derivative and excludes $W = -1$ on the relevant open domain [5, 3].

This is already good mathematical architecture. The proposed improvement is to make the polynomial and differential calculation a reusable certified computation which supplies the hypotheses of that architecture. The example thus tests the hybrid interface without requiring a different proof of the final result.

8.2 A polynomial family with an executable representation

Define $P_n \in \mathbb{Z}[w]$ by

$$P_0(w) = 1, \quad P_{n+1}(w) = (1+w)P'_n(w) - ((n+1)w + 3n+2)P_n(w). \quad (1)$$

The first values are

$$\begin{aligned} P_0(w) &= 1, \\ P_1(w) &= -w - 2, \\ P_2(w) &= 2w^2 + 8w + 9, \\ P_3(w) &= -6w^3 - 36w^2 - 79w - 64. \end{aligned}$$

The first three agree with the displayed real-polynomial identities in the inspected module; the fourth is obtained by the same recurrence. The accompanying exact-arithmetic script checks these coefficient computations. The article's argument for all indices is the recurrence and its proof, not extrapolation from these values [5].

Every operation in (1) is available over integer coefficients. A computational implementation can therefore use finite integer coefficient arrays or sparse maps and prove that their denotation follows this recurrence. Mapping coefficients to $\mathbb{R}$ yields the semantic real-polynomial family used in the calculus theorem. The inspected source instead defines its family directly as a noncomputable $\mathbb{R}[w]$ -valued object. A verified integer representation would be an implementation refinement, not a change to the mathematical recurrence.

As an example of a useful invariant, $\deg P_n = n$ and the leading coefficient is $(-1)^n n!$. Indeed, this holds for P_0 . If the leading coefficient of P_n is $c_n \neq 0$, the only degree- $n+1$ contribution in (1) is $-(n+1)wP_n$, so $c_{n+1} = -(n+1)c_n$. This proves both assertions by induction over $\mathbb{Z}$. Such invariants can be certified by ordinary theorems and then used to estimate storage or recognize a corrupted output; they need not be guessed from numerical examples.

8.3 The reusable differential computation

For $w \neq -1$, define

$$H_n(w) = \frac{e^{-(n+1)w} P_n(w)}{(1+w)^{2n+1}}, \quad \phi(w) = \frac{e^{-w}}{1+w}.$$

Put $a = n + 1$ and $k = 2n + 1$. Applying the product, chain, and quotient rules gives

$$H'_n(w) = \frac{e^{-aw}((1+w)P'_n(w) - [a(1+w) + k]P_n(w))}{(1+w)^{k+1}} \quad (2)$$

$$= \frac{e^{-(n+1)w}P_{n+1}(w)}{(1+w)^{2n+2}}. \quad (3)$$

The second equality uses

$$(n+1)(1+w) + (2n+1) = (n+1)w + 3n+2$$

and the defining recurrence. Multiplication by $\phi(w)$ now yields

$$H'_n(w)\phi(w) = H_{n+1}(w). \quad (4)$$

The exponential combination is justified by the addition law for $\exp$; all cancellations use $1+w \neq 0$. No assumption that $w > -1$ is required for this calculation. That matters because the lower Lambert branch takes values below -1 .

A certified derivative algorithm can establish (2) with P an arbitrary polynomial and n a symbolic natural-number parameter. It need not expand P_n , and it must not prove only a fixed finite collection of cases. The polynomial cleanup is a separate identity, and (4) is a separate recurrence contract. This is the useful granularity: one generic calculation schema feeding many analytic proofs.

8.4 From the computation to every iterated derivative

Theorem 8.1 (Autonomous Lambert-form derivative family). *Let $U \subseteq \mathbb{R}$ be open and $W : \mathbb{R} \rightarrow \mathbb{R}$. Assume that for every $x \in U$, W has derivative $\phi(W(x))$ at x , and $W(x) \neq -1$. With P_n defined by (1), for every $n \in \mathbb{N}$ and $x \in U$,*

$$W^{(n+1)}(x) = \frac{e^{-(n+1)W(x)}P_n(W(x))}{(1+W(x))^{2n+1}}. \quad (5)$$

Proof. Define $G_0(w) = w$ and $G_{n+1}(w) = H_n(w)$. On $w \neq -1$, the base derivative $G'_0(w) = 1$ gives $G_1(w) = G'_0(w)\phi(w)$, while (4) gives $G_{r+1}(w) = G'_r(w)\phi(w)$ for every $r \geq 1$.

We prove $W^{(r)} = G_r \circ W$ on U by induction on r . The case $r = 0$ is immediate. Fix $x \in U$ for the successor step. Since U is open, the induction identity holds on a neighborhood of x , not merely at x . The derivative of $W^{(r)}$ at x therefore agrees with that of $G_r \circ W$. The function G_r is differentiable at $W(x)$ because $W(x) \neq -1$, and the chain rule gives

$$(G_r \circ W)'(x) = G'_r(W(x))W'(x) = G'_r(W(x))\phi(W(x)) = G_{r+1}(W(x)).$$

This supplies the successor derivative and the identity. Taking $r = n + 1$ proves (5). Only differentiability of G_r at the values reached by W is used. $\square$

This is the same local mechanism captured in ProveIt's autonomous theorem. Its Lean proof explicitly turns equality on the open set into an eventual equality at the neighborhood filter before using derivative congruence [3]. Meridian should expose that mathematical premise while reconstructing the filter plumbing. Equality at a single point would be insufficient: $f(t) = t$ and $g(t) = 0$ agree at zero but have different derivatives there.

8.5 The two real Lambert branches

The inspected module applies the generic mechanism to the principal branch on $(-e^{-1}, \infty)$, where its values are greater than -1 , and the lower branch on $(-e^{-1}, 0)$, where its values are less than -1 . It provides the same polynomial formula for both, using their respective first-derivative theorems and non-pole facts [5].

This is exactly how the hybrid language should work: a single certified symbolic calculation, two separately checked analytic instantiations. A command that simply says “the other branch is similar” must reinstantiate the domain and branch hypotheses, not copy a proof receipt whose input was the principal branch. Neither endpoint $-e^{-1}$ nor the lower branch’s endpoint 0 can be added to these open-domain derivative statements by simplifying their printed formulas.

8.6 What the author would write

Here is an illustrative, unimplemented surface. The displayed recurrence is a mathematical choice, even when a CAS originally suggested it.

```
define P : Nat -> Polynomial(Int):
  P(0) := 1
  P(n+1) := (1+X)*differentiate(P(n))
            - ((n+1)*X + 3*n+2)*P(n)

let H(n,w) := exp(-(n+1)*w)*evaluate(P(n),w)/(1+w)^(2*n+1)

have step(n,w) for w != -1:
  derivative(H(n),w) * exp(-w)/(1+w) = H(n+1,w)
  by symbolic_differentiation using recurrence(P)

have all_derivatives on U:
  D^(n+1)(W)(x) = H(n,W(x))
  by autonomous_iteration using ode(W), nonpole(W), step
```

The source shown omits the expanded theorem header for readability, not from the formal declaration. The real input retains U open, its membership binders, the exact derivative hypothesis, and all quantifiers. A generated statement view shows them before publication. The symbolic method is not an oracle for the last line: the last line applies the existing theorem interface.

Obligation	Exact mathematical content	Intended provider
Polynomial denotation	Executable P_n maps to the semantic recurrence	Verified polynomial operations and an induction theorem.
Differential rule	Equation (2) for symbolic n, P, w	Derivative compiler with a checked rule for each syntax constructor.
Pole guard	$w \neq -1 \Rightarrow (1+w)^{2n+1} \neq 0$	Local arithmetic and power lemmas.
Polynomial cleanup	The numerator equals $P_{n+1}(w)$	Recurrence theorem and proof-producing polynomial normalization.
Exponential cleanup	$e^{-(n+1)w} e^{-w} = e^{-(n+2)w}$	Exponential addition and scalar algebra.

Obligation	Exact mathematical content	Intended provider
Locality	The induction identity holds near every $x \in U$	Open-set and eventual equality theorem interface.
Branch instantiation	$W(x) \neq -1$ and the ODE on the chosen domain	Principal- or lower-branch theorems, selected explicitly.

8.7 What would count as an improvement

The new contribution is not the recurrence itself, which is already present in the repository. It is a generic typed computation that supplies the derivative and normalization certificates while retaining the local analytic interface. The fair baseline receives that same computation as an ordinary Lean theorem or tactic. If the new surface adds no further benefit, the project still obtains a useful verified differential-algebra package.

Three mutations make the test discriminating. Replacing $3n + 2$ by $3n + 1$ must fail the symbolic numerator identity. Removing the non-pole hypothesis must leave a located analytic requirement, not cause the theorem header to strengthen itself. Replacing local equality by equality only at x must block derivative transfer. A system that rejects these but cannot handle the unmodified generic calculation has not yet demonstrated a useful method.

9 Finite algebra without accidental loss of generality

9.1 The kernel calculation

For $n, j \in \mathbb{N}$, consider the integer kernel

$$K_{nj} = \sum_{k=j}^n (-1)^{n-k} \binom{n}{k} \binom{k}{j}.$$

If $n < j$, the index set is empty and $K_{nj} = 0$. Otherwise put $m = n - j$ and substitute $k = j + i$, with $0 \leq i \leq m$. The elementary counting identity

$$\binom{n}{j+i} \binom{j+i}{j} = \binom{n}{j} \binom{n-j}{i}$$

is obtained by counting a j -element set contained in a $(j+i)$ -element set in two orders. Hence

$$K_{nj} = \binom{n}{j} \sum_{i=0}^m (-1)^{m-i} \binom{m}{i} = \begin{cases} 1, & n = j, \\ 0, & n \neq j. \end{cases}$$

The alternating sum is $(1-1)^m$, with the $m = 0$ case equal to 1. This reproduces the mathematical calculation exposed by the inspected orthogonality theorem [4].

A computation package need not enumerate the indices for concrete n . It can certify a symbolic index map, the counting identity, and the alternating row theorem. The demands on its certificate are exact: membership in both index sets, injectivity, surjectivity, and correspondence of summands. The inclusive upper endpoint is part of the object.

9.2 The action on a general additive commutative group

Let M be an additive commutative group and $a : \mathbb{N} \rightarrow M$. Define

$$b_n = \sum_{j=0}^n \binom{n}{j} \cdot a_j,$$

where $\cdot$ denotes integer scalar action. Rearranging a finite sum gives

$$\sum_{k=0}^n (-1)^{n-k} \binom{n}{k} \cdot b_k = \sum_{j=0}^n \left(\sum_{k=j}^n (-1)^{n-k} \binom{n}{k} \binom{k}{j} \right) \cdot a_j = a_n.$$

Only finite distributivity and the integer kernel identity are needed. No multiplication of elements of M , division in M , or embedding of M into a field occurs. This is the generality of ProveIt's module-valued result, and it must survive automation [4].

The lesson for CAS integration is to compute the *coefficients*, not to force every semantic value into the CAS coefficient domain. Mathlib's existing `match_scalars` and `module` tactics already provide relevant coefficient-wise reasoning interfaces [13]. Meridian should investigate and reuse them before building an overlapping normalization engine.

9.3 An interface that exposes the mathematical choice

```
have kernel(n,j):
  sum k in [j..n], (-1)^(n-k)*choose(n,k)*choose(k,j)
  = indicator(n=j)
cases n < j:
  conclude by empty_sum
otherwise:
  reindex k := j+i with i in [0..n-j]
  conclude by choose_product and alternating_row

conclude inversion for a : Nat -> M
by triangular_kernel_inverse using kernel
```

Here the bracket notation is explicitly for finite inclusive natural intervals. The coefficient expression is resolved in $\mathbb{Z}$; M remains an additive commutative group. The reindexing map is a visible author input. Its elementary bounds can be reconstructed under that contract.

Changing the new interval to $[0..m)$ removes a term and must invalidate coverage. Trying to replace the finite sum by an infinite sum must require a different contract, including any summability assumptions. Replacing the additive group by a semiring is also not a harmless elaboration choice: negative coefficients require the appropriate interpretation. The interface should make such mathematical differences obvious without displaying every cast lemma.

The repository's final theorem application is already concise. It should remain a negative control: a polished document view may help a reader, but there is no reason to force its author to expand a short reusable Lean expression into a longer scripted narrative.

10 Formal series, radicals, and numerical enclosures

10.1 A finite jet is not an infinite solution

Work in formal power series over $\mathbb{Z}$ and consider

$$C(X) = 1 + XC(X)^2. \tag{6}$$

A CAS can compute the jet

$$c(X) = 1 + X + 2X^2 + 5X^3 + 14X^4 + 42X^5 \pmod{X^6}.$$

An exact polynomial calculation verifies that $c - 1 - Xc^2$ has zero coefficients below degree six. Its coefficient of X^6 is -132 , so c is not itself an exact polynomial solution of (6). The correct certificate is equality in a truncated series ring, not equality in $\mathbb{Z}[[X]]$.

There is nevertheless a complete formal construction. Define

$$c_0 = 1, \quad c_n = \sum_{i=0}^{n-1} c_i c_{n-1-i} \quad (n \geq 1), \quad C = \sum_{n \geq 0} c_n X^n.$$

The coefficient of X^0 in (6) is $c_0 = 1$. For $n \geq 1$, the coefficient of X^n on the right is precisely the displayed finite sum. Thus the recursively defined series satisfies the equation. Conversely, every solution has the same constant coefficient, and strong induction forces every successive coefficient by that formula. This proves existence and uniqueness as a formal series. The proof uses only ring operations and finite coefficient sums; it does not invoke analytic convergence.

Meridian should represent these as different products of computation:

```
calculate c := jet_solution(C = 1 + X*C^2, order = 6)
# Contract: the residual is zero modulo X^6.

construct C : FormalSeries(Int)
  satisfying C = 1 + X*C^2 uniquely
  by coefficient_recursion
# Contract: existence, the full identity, and uniqueness.
```

The corresponding analytic expression involving $\sqrt{1-4x}$ would require additional work: a real or complex branch, a domain, a treatment of the removable value at zero, and a convergence or analytic identity theorem. A coefficient identity cannot discharge those obligations by changing its label.

A useful further refinement is an order-aware jet interface. To compute a product modulo X^N , each input need only be known modulo X^N . Differentiation usually requires one more input order to guarantee the requested output order. A substitution has its own admissibility and order contract. These are ideal verified algorithms because the dependency on the requested precision is mathematical and reusable.

10.2 A radical candidate needs branch evidence

Consider the denesting identity

$$\sqrt{5 + 2\sqrt{6}} = \sqrt{2} + \sqrt{3}.$$

A CAS can propose the right-hand side. A rigorous checker verifies its square:

$$(\sqrt{2} + \sqrt{3})^2 = 2 + 3 + 2\sqrt{2}\sqrt{3} = 5 + 2\sqrt{6},$$

using the nonnegativity of the radicands for the square-root product rule. It also checks $\sqrt{2} + \sqrt{3} \geq 0$. The nonnegative square root is unique, so the identity follows. The negative of the candidate has the same square and would fail the branch condition.

This is a useful small external-CAS certificate interface: return an algebraic candidate, verify its equation, and verify the selected real embedding or sign. For more complicated algebraic numbers, an isolating interval can identify the chosen root. A minimal polynomial alone does not select an embedding. A denesting algorithm's failure to find a candidate is not proof that no expression in a proposed radical class exists; such a nonexistence claim needs a precisely defined class and an appropriate completeness theorem.

The universal identity $\sqrt{x^2} = |x|$ supplies a nearby test. Replacing the right side by x adds the requirement $x \geq 0$; at $x = -1$ it is false. The normalizer must retain the absolute value or expose the condition, not silently take the branch that makes the expression shorter.

10.3 Exact bounds are different from approximate values

For a numerical-looking computation, the output contract can be a rational interval. For instance,

$$\frac{1414213}{10^6} < \sqrt{2} < \frac{1414214}{10^6}$$

follows from nonnegativity and the exact integer inequalities

$$1999998409369 < 2000000000000 < 2000001237796.$$

These are the squared numerators against $2 \cdot 10^{12}$. A floating-point routine may discover the interval, but the interval's authority comes from the exact inequalities and monotonicity of squaring on nonnegative reals. The accompanying script checks these integers without floating-point arithmetic.

For a complicated integral, an interval certificate must additionally account for the integrand's domain, the integration convention, and any truncation or rounding bounds. A small residual from numerical quadrature is not such a certificate. Existing verified enclosure systems provide important precedents, but the exact imported theorem remains the acceptance boundary [18].

10.4 Summation and asymptotics need their own result languages

An indefinite finite-sum antidifference $F(k+1) - F(k) = f(k)$ is a useful certificate, but using it over a range requires the endpoint evaluation. A recurrence for a sequence requires initial conditions to identify a particular sequence. A closed form involving special functions requires a domain and convention for those functions. These conditions should be attached to the computed object.

Similarly, an asymptotic expansion is neither a formal series identity nor a numerical approximation at a few sample points. A contract might assert, for example,

$$\exists M > 0 \exists x_0 \forall x \geq x_0, \left| f(x) - \sum_{k=0}^{N-1} a_k \psi_k(x) \right| \leq M |\psi_N(x)|.$$

The asymptotic scale, the approach filter, the domain, and whether M, x_0 depend on N are mathematical inputs. A symbolic engine can propose coefficients; a theorem must supply the remainder contract. This is a natural future extension of the same interface, not a reason to interpret every CAS series command as already proving an asymptotic theorem.

11 Leant as a partner in the hybrid system

11.1 Keep structural synthesis distinct from symbolic mathematics

Leant already supplies a useful architecture: a narrow synthesis boundary, candidate groups, retained origins, configurable search, and a verification stage. The inspected engine describes its output as candidate batches rather than as behavioral authority; its verification module deliberately treats its opaque receipt as acceptance by a callback [7, 6].

Meridian should preserve that separation. Leant can assemble implications, instantiate available providers, and connect the hypotheses produced by one computation to the next theorem. It need not become the polynomial factorizer or the engine which chooses a branch of an inverse

function. Conversely, a CAS should not be asked to decide the validity of a general dependent proof context. The typed workspace supplies a common boundary while retaining specialized providers.

For example, a derivative compiler returns a theorem about H'_n ; a polynomial normalizer returns the recurrence identity; a local theorem supplies $W(x) \neq -1$. Leant can help assemble the required theorem application. Its candidate must still be checked against the exact instantiation of the analytical contract, including the range-local hypotheses.

11.2 A common request format

A provider request should have the following logical content; the wire format may use compact references to Lean-owned expressions.

```
Request:
  request_id, environment_id, scope_id
  ordered_telescope, exact_target_or_result_contract
  accepted_input_object_ids
  permitted_local_premises, permitted_library_inventory
  axiom_policy, method_requirement
  result_representation, domain_and_branch_policy
  global_resource_budget, source_origin
```

The external provider is not given mutable authority over the target. A typed input object has a semantic identity independent of how it prints. An opaque atom may be exported as an identifier referring back to a Lean-owned expression; its printed spelling is explanatory metadata, not identity.

A CAS-facing request can additionally specify a coefficient ring, variable order, and accepted expression grammar. A Leant-facing request can contain a structural projection and a map back to the exact source context. Both return through the same exact acceptance gate. Neither provider's internal abstraction is allowed to replace the original problem.

11.3 Outcomes and negative evidence

The operational response variants should distinguish a candidate, a proposed refinement with child obligations and a typed combiner, an unsupported fragment, a budget exhaustion, a rejected candidate, and an accepted result with retained evidence. A checked counterexample or a proof of nonexistence is another positive piece of evidence, not the absence of a successful search.

Leant's engine comments use a two-part condition for strong negative verdicts: the engine's claimed completeness and the absence of unsafe abstraction on the Leant side [7]. That is a valuable discipline for labeling search outcomes. It is not, by itself, a replacement for the required logical bridge. For publication, a nonexistence assertion needs a proof of the exact original negative proposition, or a checked certificate and a proved adequacy theorem which entail it. An empty list of abstraction losses does not establish an adequacy theorem merely by being empty.

For a proposition P , the strong negative result is evidence of $\neg P$. For a data specification $\text{Spec}(x, y)$, it is evidence of $\neg \exists y, \text{Spec}(x, y)$. A budgeted solver returning no candidates proves neither. A language model may propose an actual counterexample and submit its verification like any other provider; its confidence or failed search is not a negative certificate.

11.4 A suggestion is not the same as its displayed replacement

The committed `Main.hs` inspection confirms an important distinction from the syntheses. The suggestion path probes a tactic against a persistent proof state, checks for errors and a successor

state, and can display extracted `Try this` text in place of the tactic that ran. The closure annotation uses a reduction in the visible goal count together with a textual check for a remaining-subgoals comment. The applied-tactic path can also store the extracted text alongside the state produced by the original tactic [8].

These mechanisms are reasonable draft-REPL behavior; they should not be misdescribed as a proof that the exact displayed edit independently replays to a completed document. Meridian needs four separate statuses:

Status	Meaning
Probe accepted	This attempted operation produced a valid successor according to the backend protocol.
Displayed edit replayed	The exact proposed replacement was executed from the same immutable origin and independently accepted.
Requested result checked	The exact claim or output specification has retained logical evidence under the policy.
Document sealed	Every asserted node, including unused claims, and the final target pass publication checks.

A decrease in goal count can be useful progress without proving the whole request. Missing or malformed status fields must be protocol failures, not interpreted as proof completion. Before caching a “verified suggestion,” the adapter should replay the exact displayed replacement and bind its receipt to the request, environment, origin state, and accepted candidate.

11.5 Transactional search and accepted data

Two obligations may share a synthesized polynomial, a witness, a typeclass instance, or a metavariable. Solving them independently and keeping incompatible assignments is not a valid refinement. Provider attempts therefore occur on an immutable snapshot; dependent choices are committed together after their joint constraints are checked. This applies equally to Leant and to CAS output that creates auxiliary objects.

Proof-only search may select any admissible proof of the fixed target. Data search is different. If two functions have the same type but different behavior, they remain different candidates. Even two values satisfying a weak common specification can affect later computations differently. Changing the accepted value is an object change unless a registered observational equivalence, together with all needed respectfulness theorems, authorizes the replacement. The first implementation should simply retain the chosen data value rather than attempt universal observational equivalence.

Resource limits should cover a whole request, not just each candidate. A sequence of individually bounded tactic probes can otherwise consume unbounded total work. Process failure, malformed responses, and cancellation are operational outcomes; none establish a mathematical result. Separate worker processes can protect an interactive session, while the fresh validator protects the evidence boundary. These are different purposes.

12 A realizable architecture and evidence format

12.1 One protocol, two authoring frontends

The initial implementation should remain a Lean package. A new global parser, an independent elaborator for all mathematics, and a new logical kernel are unnecessary. Structured steps can




Figure 1: Proposed architecture. Provider success is not the publication gate; validation uses the original typed request. The mathematical source frontend is optional, and shares its computational infrastructure with ordinary Lean.

first be expressed through Lean syntax extensions and a generated document view. This makes existing source usable before a separate language becomes a requirement.

The workspace should own immutable elaborated inputs. It should not repeatedly reparse a pretty-printed telescope for every provider. Compact Lean-owned identifiers can reference local expressions while an exported request includes enough explicit binding structure for independent reconstruction. Closed proof obligations are formed by abstracting their actual ordered contexts, not by sorting hypotheses alphabetically.

12.2 The logical acceptance rule is intentionally modest

A complete document is accepted only when its final proof has the fixed target, its declared computed objects have their specified properties, and every assertion that the document presents as proved has evidence. A proof of the final theorem which ignores a false intermediate assertion does not rescue the published document. Unused but true assertions may be marked as unused.

Recursive mathematical definitions and induction remain available through Lean’s checked recursion and induction principles. The workspace must not confuse an admissible recursive definition with an unresolved cycle of proof debts.

The familiar relative soundness argument then follows from checked theorem application, dependent substitution, and a final check. There is little value in presenting that argument as a new foundational theorem. The new proof obligations in Meridian concern the concrete computational specifications, representation bridges, domain preservation, and correctness of the algorithms or checkers. Those are the theorems whose implementation would make the hybrid useful.

Logical validity also does not guarantee statement fidelity or explanatory quality. Lean’s validation guidance explicitly distinguishes the validity of a proof from the meaning of the formal statement [9]. Meridian keeps both a resolved semantic header and a human-readable rendering, but does not claim that either removes the need for mathematical review.

12.3 Three retained forms of evidence

An accepted computation should retain three complementary forms. The *mathematical record* contains its original inputs, accepted output, domain, exact relation, branch policy, and method. The *certificate record* contains the exact data consumed by a checker, the checker identity and version, and the theorem relating checker acceptance to the specification. The *host evidence* contains the accepted proof term or sufficient pinned proof data for its revalidation. Readable Lean source remains a companion form.

Search-free replay means that discovery providers are not consulted. It does *not* mean computation-free checking: kernel reduction or deterministic certificate checking may still perform substantial work. A term containing a reflected checker can require that checker to reduce again. The artifact should report this cost rather than imply that retaining a term makes all verification instantaneous.

For strict term replay, the validator accepts no arbitrary tactic source from the producer; it checks supplied declarations against a controlled environment. For certificate replay, it runs a pinned, deterministic certificate interpreter which constructs evidence and then checks that evidence. A saved tactic script which invokes open search is a third, weaker reproducibility artifact and must not be mislabeled as search-free replay.

12.4 A concrete artifact layout

```
proof-project/
  document.lean           # Authoring source or Lean companion
  document.meridian       # Optional structured source
  semantics.json          # Typed identities and interpretation notices
  objects/                # Exact accepted computational data
  certificates/           # Inputs to pinned checkers
  evidence/               # Host proof evidence and declaration identities
  source-map.json         # Claims, computations, and visible explanations
  environment.lock        # Toolchain, imports, definitions, packages
  audit.json              # Axioms, outcomes, conformance, replay mode
```

This is a proposed runtime format, not a description of files generated by the present article. Its exact serialization is an implementation task. A hash provides identity and tamper detection relative to trusted metadata; it is not proof. Imported binary modules also require a trust policy: an opaque compiled file is not an independently checked theorem merely because its filename is recorded.

When a library is upgraded, the old artifact remains a claim about its old pinned environment. Migration must first compare the semantic declarations and then recheck evidence under the new environment. If old terms and newly elaborated source disagree, neither is automatically authoritative for the intended new statement. The migration stops for review. Proof changes may be automatic within policy; plan, object, and statement changes receive distinct visible diffs, as synthesis A recommends [1].

12.5 Operational safety and privacy

An external provider can be mathematically untrusted and still execute dangerous code. Its process therefore receives only the permissions needed for the request, and a remote CAS does not receive private definitions or full contexts without an explicit project policy. Read-only filesystem access, bounded resources, and no implicit network access are sensible defaults for a local computation worker.

These engineering protections are not new logical axioms. They protect the user’s environment and interactive state. The validator must also obtain the intended target from a source outside the producer’s authority, so that a producer cannot make a false result appear checked by replacing the target with `True`. The operational boundary and the statement boundary solve different problems and both belong in a deployed system.

13 A bounded implementation plan

13.1 First slice: one honest computed object

The first end-to-end slice should be domain-preserving rational simplification with complete solving of small univariate examples. It is small enough to specify fully and rich enough to test the central contribution. It includes a partial-expression package, a proof-producing polynomial normalizer, exact input reification, retained domain data, a result contract, and a document view.

The acceptance gate is concrete: the quotient in Section 6 simplifies while retaining its missing point; the complete solution set depends on the deliberately selected partial or total convention; specialization at the missing point fails in partial mode. The same operations are available from ordinary Lean. A broad language with a weak “verified” badge is not an acceptable substitute for this slice.

13.2 Second slice: a parametric differential certificate

Implement a small derivative compiler and a polynomial representation package, then produce the parameterized derivative calculation in Section 8. It must handle arbitrary polynomial input and symbolic n , not only output a list of the first few derivatives. The certificate is connected to ProveIt’s actual autonomous theorem under its existing assumptions.

The gate is a replayable proof of the generic step and both branch instantiations, with no external search during replay. The negative neighbors alter the recurrence, remove the non-pole hypothesis, and weaken neighborhood equality to pointwise equality. This gives analysis an early, concrete role, consistent with the syntheses’ concern that analysis plumbing was underrepresented in some samples [1, 2].

13.3 Third slice: structural synthesis and independent suggestions

Connect Leant through the common typed protocol after measuring it on the structural obligations produced by these two slices. Compare it with direct lemma application and existing Lean search supplied the same provider inventory. Repair the displayed-edit replay boundary before advertising independently verified suggestions.

The gate is interchangeability: disabling Leant must not invalidate a stored proof; replacing it with another candidate producer must not change acceptance. A provider timeout, malformed result, or incorrect output should fail locally without corrupting the current accepted mathematical objects.

13.4 Fourth slice: representations and finite algebra

Add the three-view stress test and the binomial-transform example. Verify that integer coefficient computation acts on arbitrary additive commutative groups without importing a field requirement. Then add formal jets, with order propagation and coefficient projection contracts.

The gate is not a fixed promised compression percentage. It is evidence that these packages remove recurring author work, expose better failure messages, and preserve theorem generality under controlled edits. Helpers and package maintenance are charged equally to the ordinary-Lean

comparison.

13.5 What should be deliberately deferred

A general natural-language definition parser, universal simplification of transcendental expressions, arbitrary certified symbolic integration, global conditional equality saturation, cross-foundation proof transport, and automatic observational replacement of arbitrary witnesses should not be initial requirements. None is needed to test the proposed central interface.

The most dangerous implementation trap is to build many attractive verbs whose actual contract is “try a large tactic and print success.” A smaller family of operations with exact data, domain, and result guarantees would be a more useful second-round outcome.

14 Evaluation: isolate the language from the mathematics

14.1 Comparison conditions

An evaluation should separate original code, improved mathematical libraries, improved algorithms, and the new authoring surface. At minimum, compare:

Arm	Configuration
A	Original ProveIt Lean as a historical reference, not the sole baseline.
B	Refactored Lean with the same theorem wrappers, verified algorithms, certificate checkers, and allowed search inventory as Meridian.
C	Arm B with a generated mathematical document view and retained computation records, but ordinary Lean authoring.
D	The proposed mathematical source frontend over exactly the same backend as B and C.

The difference between A and B estimates the combined library and automation improvement, not a language advantage. The difference between B and C tests the value of presentation and accounting. The difference between C and D tests whether the new input form is worth learning. A second ablation should enable or disable the new CAS method while holding the frontend fixed.

The target theorem, aliases that expose it, and its downstream closure must be excluded from discovery. A search service that simply retrieves the theorem being reconstructed invalidates the comparison. Cached target-specific results must be handled by the same exclusion policy. Split by theorem families or modules, rather than random neighboring proof lines, to reduce leakage through near-identical lemmas.

14.2 Measure authoring, reading, computation, and maintenance separately

Authoring measures should include time to a checked result, number of explicit mathematical choices, number of attempts, manual refinement of guards, and time spent diagnosing failure. Reader tasks should ask participants to identify the main invariant, explain why a domain condition is necessary, and predict the failure caused by a nearby mutation. Agreement with a stylistic preference survey alone would be weak evidence.

Computation measures include elaboration time, search time, certificate generation, certificate size, and replay time, with failures and high-percentile costs reported. Maintenance measures include changes needed after a theorem wrapper, instance, representation route, or library revision

changes. The cost of writing and maintaining a domain package must be amortized over its actual uses, not hidden outside the accounting.

A useful error measure is *contract surprise*: the fraction of operations whose displayed guarantees a reader overestimates. Asking whether a returned factorization is irreducible, whether a jet is an infinite solution, and whether a missing branch is covered directly tests the intended value of the interface. Another is *guard surprise*: whether the author can predict which specializations remain legal after a calculation.

14.3 Negative neighbors are executable specifications

The following are proposed acceptance tests for an implementation, not claims that a Meridian implementation already passed them. The companion Python script checks selected arithmetic examples only.

ID	Mutation or failure	Required behavior
D1	Cancel $(x^2 - 1)/(x - 1)$ and evaluate the partial result at 1	Retain the excluded point; reject the evaluation.
D2	Use a partial-function solution certificate for the totalized equation	Reject the semantic mismatch; the complete solution sets differ.
D3	Solve $ax = 0$ as $\{0\}$ without considering $a = 0$	Report a conditional answer or require branch coverage.
D4	Substitute a value at which an inherited guard fails	Reinstantiate the guard and locate its origin; do not use the cached equality unconditionally.
C1	Verify only that returned factors multiply to the input	Do not label them irreducible without further evidence.
C2	Verify every proposed root, but omit one	Accept only root soundness; reject a complete-answer claim.
C3	Supply a tiny floating residual for an exact identity	Treat it as a candidate diagnostic, not proof.
C4	Use a Gröbner membership witness with rational coefficients over $\mathbb{Z}$	Reject the coefficient-domain mismatch.
V1	Cast $a - b$ to $a - b$ without $b \leq a$	Keep the guard or select an explicit alternative representation.
V2	Add a competing representation route	Preserve the old route or request a visible selection; never choose by provability.
V3	Cancel a nonzero nonunit in a ring with zero divisors	Require the actual cancellation capability.
A1	Change $3n + 2$ to $3n + 1$ in the Lambert recurrence	Reject the symbolic numerator identity.
A2	Remove $W(x) \neq -1$ from the analytic instantiation	Expose the unresolved requirement without strengthening the target.
A3	Differentiate from equality at one point only	Require neighborhood equality or a different valid theorem.
F1	Omit the upper endpoint in a finite reindexing	Fail coverage.

ID	Mutation or failure	Required behavior
F2	Treat a finite jet certificate as a full formal identity	Reject the relation-strength mismatch.
L1	Provider returns a different expression with the same printed name	Reject incorrect typed origin binding.
L2	Displayed replacement differs from the probed tactic	Independently replay the exact replacement before giving it a replay receipt.
L3	A request times out or a status field is missing	Return exhaustion or a protocol error, never completion or refutation.
L4	Two children commit incompatible shared data	Reject the combined transaction.
P1	A false intermediate claim is unused by the final proof	Reject document completion despite the final theorem's proof.
P2	A transitive dependency uses an unapproved axiom	Fail the declared publication policy.
P3	A library update changes the resolved statement but preserves its spelling	Invalidate migration and display a semantic change.

14.4 Stopping rules

If the computational packages are useful but the new syntax does not improve authoring or comprehension, ship the packages and the document view. If explicit guards create more confusion than they prevent, improve their grouping and rendering rather than deleting their semantics. If a verified algorithm is too expensive for interactive use, retain it as a reference and investigate a certificate-producing external implementation with the same contract.

Success is not the existence of a language bearing a new name. It is a measured reduction in the cost of developing, checking, understanding, and maintaining the intended mathematics, without relocating that cost into unreported infrastructure or unexamined assumptions.

15 Decisions for the second brainstorm iteration

The following answers consolidate the concrete questions in both syntheses. They are design choices to test, rather than claims that the associated engineering has already been completed [1, 2].

Open question	Proposed decision and discriminating evidence
Is the language an input form or a document?	Both frontends share the same workspace. Compare a generated view with authored structured source before making the latter mandatory.
What is the smallest useful method?	A project-owned theorem interface with role annotations for logical steps; a result contract plus a verified algorithm or checker for computation. Demonstrate one end-to-end rational example first.
Must the displayed reason constrain the proof?	Yes when it states a particular method. Keep an explicitly unconstrained search form. Test that a correct but different proof triggers a reason change.

Open question	Proposed decision and discriminating evidence
Which choices may be automatic?	Proof evidence within a fixed contract may be replaced under policy. New accepted data, branches, plans, and statements are visible changes.
How are view chains composed?	Retain the actual intermediate objects and instantiate every guard along the route. Begin with equality-based denotations; require explicit routes or checked coherence.
Where does definition authoring start?	With partial-expression and polynomial representation packages. Compare downstream work against alternate ordinary Lean representations under the same specification.
How are changing theorem binders handled?	Register roles on maintained semantic wrappers. Recheck wrapper types and test migration; do not infer meaning from binder position.
How are semantic notices dismissed?	Record a dismissal against the resolved construction and selected convention. Invalidate it when that meaning changes; measure false alarms. A dismissed notice never discharges a method obligation.
What does verified suggestion mean?	Distinguish probe acceptance, exact-edit replay, result certification, and document completion. Test malformed and partial responses.
What is the provider's negative interface?	Return search status and abstraction loss metadata, but require actual evidence and a proved bridge for mathematical nonexistence.
Is Leant worth an adapter?	Measure its contribution on multi-step structural assembly against equally equipped Lean search. Do not assume every <code>exact</code> line is a synthesis opportunity.
What is the unit of replay?	A typed mathematical node plus its accepted data and evidence under a pinned environment. Terms and portable certificates have explicitly different replay modes.
Must discarded data candidates be published?	No. Publish the accepted value and its specification; retain optional exploration logs separately under privacy policy.
Who owns competing method packages?	Semantic interfaces have owners and explicit versions. Conflicting interpretations need qualification, not import-order precedence.
How is infrastructure cost charged?	Share every wrapper, algorithm, and checker with the Lean baseline; report construction and maintenance effort.
What would justify abandoning the surface?	Evidence that the packages and view help but a new input form does not. A library-and-tooling outcome is an explicit success condition.

16 The remaining research questions

The most promising research problem is *demand-directed certification*. A consumer should state what it needs from a computed object, and the system should seek no stronger certificate than necessary. If a proof only needs a polynomial consequence, a membership identity may be enough; proving a full Gröbner basis is wasted effort. If a proof later needs completeness, the same accepted object can acquire an additional property without changing its identity. This requires a

library of proved implications between result contracts, not a heuristic confidence hierarchy.

A second problem is *computation-aware explanation*. A factorization trace may be enormous, but the useful mathematical reason may be “this product gives the two real roots, and the other factors have no real zeros.” An explanation should summarize the certified properties actually used. Method conformance can check that a displayed construction follows a specified interface, but it cannot by itself prove that the explanation is pedagogically good. Reader tasks must supply that evidence.

A third problem is *domain minimization with proofs*. A computation can start with conservative guards and later prove that some are unnecessary, or replace a conditional expression with a piecewise one covering more of the domain. These operations should have their own certificates. They are attractive because they improve mathematical usability without sacrificing exactness, but no general algorithm for weakest domain conditions is assumed.

A fourth problem is *representation choice under a stable specification*. Sparse and dense polynomial forms, algebraic-number encodings, and alternative formal-series representations can have very different computational costs. A verified refinement interface can permit implementation changes while preserving the public object. Optimizing this choice without hiding changes to accepted data or observations is a substantive compiler problem.

Finally, *proof extraction from existing developments* can make adoption incremental. Existing **have** chains, calculations, and recognizable computational fragments can be lifted into workspace nodes. The extracted view should initially preserve the exact proof rather than invent an alleged author intention. User edits can then group or name meaningful steps. This would also provide a realistic corpus for evaluating whether the proposed units match mathematicians’ actual working patterns.

17 Conclusion

The first iteration found a sound foundation for a readable Lean-facing proof layer. The most useful second step is not another set of near-synonymous proof verbs. It is to make *certified mathematical computation* a first-class part of that layer.

Meridian would let the author select an object, a domain, a representation, or an ansatz, then ask for computational consequences with explicit guarantees. A verified algorithm can produce a result; an external CAS can suggest one; Leant can assemble the logical connections. None receives authority to alter the intended statement or inflate the strength of a returned certificate.

The Lambert example shows the intended division of labor. A reusable symbolic calculation establishes the polynomial differential recurrence. A local analytic theorem transports that calculation through a solution of an autonomous equation. Branch hypotheses remain visible, and the same computation serves both branches. The rational-expression example shows why the interface must also preserve meaning: dropping a domain can change a complete solution set even when a familiar cancellation is valid wherever it was originally defined.

The proposed system is ambitious in its mathematical interface but conservative in its foundations and initial scope. Its first deliverables should be small verified computational packages, exact result contracts, and readable retained evidence. A new authoring language is justified only to the extent that it makes those capabilities easier to use and the resulting mathematics easier to understand.

A Minimal protocol and semantic interfaces

The following Lean-like type sketches describe interfaces; they are not compiled source in this deliverable. Universe parameters, serialization, and operational metadata would be supplied by an implementation.

```

structure Certified (A : Type u) (P : A -> Prop) where
  value : A
  evidence : P value

structure ConditionalResult (A : Type u) (P : A -> Prop) where
  guard : Prop
  value : A
  evidence : guard -> P value

structure CheckedAlgorithm (I : Type u) (O : I -> Type v)
  (spec : (x : I) -> O x -> Prop) where
  run : (x : I) -> O x
  correct : (x : I) -> spec x (run x)

structure CertificateInterface (I : Type u) (O : I -> Type v)
  (C : (x : I) -> O x -> Type w)
  (spec : (x : I) -> O x -> Prop) where
  check : (x : I) -> (y : O x) -> C x y -> Bool
  sound : (x : I) -> (y : O x) -> (c : C x y) ->
    check x y c = true -> spec x y

```

The constant-value conditional interface is adequate when the returned symbolic representation exists independently of its guard. For an object whose type or construction itself depends on a guard witness, use a dependent function or scoped branch instead; the displayed simplified interface must not erase that dependency.

A provider does not create a trusted **Certified** value by emitting JSON with an **evidence** field. The host parses the returned data, checks the relation to the original input, reconstructs the alleged evidence, and checks its type and dependencies. Callback receipts and logical evidence are distinct objects.

A minimal proposed surface grammar is:

```

document      ::= declaration*
declaration   ::= context | definition | theorem | computation
context       ::= "context" name ":" scoped_declaration*
computation   ::= "calculate" name ":" request contract_clause*
step          ::= claim | obtain | computation | cases | induction
               | calculation_chain | lean_escape
claim         ::= "have" name ":" proposition justification
contract_clause ::= "on" domain | "preserving domain"
               | "completely" | "modulo" precision
               | "satisfying" specification
justification ::= "by" method input_clause* | "by search" policy

```

Formulas initially reuse Lean's term elaboration. The grammar is intentionally small; the result contracts and the typed method registry determine what each operation means. New notation is a package feature, not a source of permission for search to reinterpret a formula.

B Source register and inspection boundaries

B.1 Pinned repository revisions

Brainstorm.

494be9d5ab41a5ae219fc6f800e6aee3c93bce39. Both requested synthesis sources were read, including their comparative assessments, Leant discussion, proposed experiments, and next-iteration questions. Individual proposals were used through those reports’ detailed accounts; this article does not claim an independent full rereading of all nine originals.

ProveIt.

24ce8bd743eaab64a91ce90725ea00f498d319d2. Direct static inspection covered the three modules listed below. Synthesis B’s corpus audit used local revision 7c4e3f10; its statistics are attributed rather than asserted as a new measurement at this review’s public ref.

Leant.

3a40904be8a410d832d9ab6900b3d3e7b425eccb. Direct static inspection covered the verification receipt boundary, the engine’s interface and module commentary, and relevant tactic-suggestion and applied-edit code in `Main.hs`. Uncommitted behavioral work discussed in the syntheses was not executed or verified here.

B.2 Inspected file locators

The ProveIt module root is `Analysis/FabiusFunction/Lean/FabiusFunction/`. Under it, the reviewed files are `AutonomousIteratedDeriv.lean`, `BinomialInversion.lean`, and `LambertWHigherDerivatives.lean`. The article’s analysis uses their actual mathematical decomposition, not an assumption that every long proof can be compressed automatically.

The Leant paths are `src/Leant/Synth/Verification.hs`, `src/Leant/Synth/Engine.hs`, and `src/Main.hs`. In the last file, the review focused on source lines 4600–4760 and 4820–5110 at the pinned revision: response parsing, displayed suggestions, bounded probes, chaining, and storing applied edits. No claim is made that the Haskell project was built or that the live backend was exercised.

The bibliography links every repository item to its pinned revision. External manuals and current documentation were consulted on September 5, 2026, Pacific time. The moving Lean reference identified itself as covering version 4.34.0-rc2; that is documentation context, not a claimed ProveIt toolchain or a tested compatibility target [12].

C Companion calculations and limits of validation

The archive includes `check_examples.py` and its recorded output `example_checks.json`. The script uses Python integers and exact rational arithmetic from the standard library. It checks polynomial product identities, the first Lambert recurrence values, a coefficient identity in a symbolic parameter, finite derivative-numerator instances, the formal-series jet residual, integer square-root bounds, and several nearby false variants.

The execution recorded **27 of 27 checks passed**. That result is a test of the supplied arithmetic examples and the script, not a formal proof of the script’s correctness. It is not an evaluation of Meridian, a Leant integration, or a Lean kernel certificate. The universal arguments about partial-expression composition, the derivative family, and the formal-series solution are proved mathematically in the article; they have not been compiled as new Lean theorems in this deliverable.

The article therefore makes three different claims with different support: observations about

repository code are based on pinned static reads; the language and protocol are proposed designs; and the arithmetic supplement was actually executed. Keeping these claims separate is itself an instance of the result contract discipline advocated throughout the article.

References

- [1] Codex. *Mathematical Ideas, Checked Evidence: A Unified Evaluation of Nine Proof-Language Proposals*. Brainstorm synthesis, September 5, 2026. <https://github.com/VladimirReshetnikov/Brainstorm/blob/494be9d5ab41a5ae219fc6f800e6aee3c93bce39/docs/synthesis/unified-report.tex>.
- [2] Claude (Anthropic). *Nine Blueprints for a Mathematician’s Proof Language over Lean*. Brainstorm synthesis, revised September 5, 2026. https://github.com/VladimirReshetnikov/Brainstorm/blob/494be9d5ab41a5ae219fc6f800e6aee3c93bce39/docs/unified_report/unified_report.tex.
- [3] Vladimir Reshetnikov and contributors. *ProveIt: AutonomousIteratedDeriv.lean*. Source inspected at revision 24ce8bd7. <https://github.com/VladimirReshetnikov/ProveIt/blob/24ce8bd743eaab64a91ce90725ea00f498d319d2/Analysis/FabiusFunction/Lean/FabiusFunction/AutonomousIteratedDeriv.lean>.
- [4] Vladimir Reshetnikov and contributors. *ProveIt: BinomialInversion.lean*. Source inspected at revision 24ce8bd7. <https://github.com/VladimirReshetnikov/ProveIt/blob/24ce8bd743eaab64a91ce90725ea00f498d319d2/Analysis/FabiusFunction/Lean/FabiusFunction/BinomialInversion.lean>.
- [5] Vladimir Reshetnikov and contributors. *ProveIt: LambertWHigherDerivatives.lean*. Source inspected at revision 24ce8bd7. <https://github.com/VladimirReshetnikov/ProveIt/blob/24ce8bd743eaab64a91ce90725ea00f498d319d2/Analysis/FabiusFunction/Lean/FabiusFunction/LambertWHigherDerivatives.lean>.
- [6] Vladimir Reshetnikov and contributors. *Leant: Verification.hs*. Source inspected at revision 3a40904b. <https://github.com/VladimirReshetnikov/Leant/blob/3a40904be8a410d832d9ab6900b3d3e7b425eccb/src/Leant/Synth/Verification.hs>.
- [7] Vladimir Reshetnikov and contributors. *Leant: Engine.hs*. Source inspected at revision 3a40904b. <https://github.com/VladimirReshetnikov/Leant/blob/3a40904be8a410d832d9ab6900b3d3e7b425eccb/src/Leant/Synth/Engine.hs>.
- [8] Vladimir Reshetnikov and contributors. *Leant: Main.hs*. Source inspected at revision 3a40904b, especially the suggestion and applied-tactic paths. <https://github.com/VladimirReshetnikov/Leant/blob/3a40904be8a410d832d9ab6900b3d3e7b425eccb/src/Main.hs>.
- [9] The Lean development team. *Validating a Lean Proof*. Lean Language Reference. <https://lean-lang.org/doc/reference/latest/ValidatingProofs/>. Accessed September 5, 2026.
- [10] The Lean development team. *Axioms*, especially “Creating and Tracking Proofs That Trust the Compiler” and “Axioms and the native_decide Tactic.” Lean Language Reference. <https://lean-lang.org/doc/reference/latest/Axioms/>. Accessed September 5, 2026.
- [11] The Lean development team. *Tactic Reference and Lean 4.30.0 release notes*, concerning call-by-value proof-layer simplification. <https://lean-lang.org/doc/reference/latest/Tactic-Proofs/Tactic-Reference/>; <https://lean-lang.org/doc/reference/latest/releases/v4.30.0/>. Accessed September 5, 2026.
- [12] The Lean development team. *Introduction*. Lean Language Reference, version identified by the consulted page as 4.34.0-rc2. <https://lean-lang.org/doc/reference/latest/Introduction/>. Accessed September 5, 2026.

- [13] Mathlib contributors. *Mathlib.Tactic.Module: match_scalars and module*. https://leanprover-community.github.io/mathlib4_docs/Mathlib/Tactic/Module.html. Accessed September 5, 2026.
- [14] Mathlib contributors. *Mathlib.Tactic.Polyrith*. Current documentation notes that the external service is no longer available and the tactic is unsupported. https://leanprover-community.github.io/mathlib4_docs/Mathlib/Tactic/Polyrith.html. Accessed September 5, 2026.
- [15] James H. Davenport. *Towards Verified Polynomial Factorisation*. arXiv:2409.09533v1, September 14, 2024. <https://arxiv.org/abs/2409.09533>.
- [16] Hao Shen, Junyu Guo, Junqi Liu, and Lihong Zhi. *Automated Tactics for Polynomial Reasoning in Lean 4*. arXiv:2604.13514v1, April 15, 2026. <https://arxiv.org/html/2604.13514v1>.
- [17] Jose Divasón, Sebastiaan J. C. Joosten, René Thiemann, and Akihisa Yamada. *The Factorization Algorithm of Berlekamp and Zassenhaus*. Archive of Formal Proofs, October 14, 2016, with subsequent revisions. https://isa-afp.org/entries/Berlekamp_Zassenhaus.html. See also *A Verified Implementation of the Berlekamp–Zassenhaus Factorization Algorithm*, Journal of Automated Reasoning 64, 699–735, <https://doi.org/10.1007/s10817-019-09526-y>.
- [18] Guillaume Melquiond and contributors. *Interval Package for Rocq*. Official documentation, accessed September 5, 2026. <https://coqinterval.gitlabpages.inria.fr/>.
- [19] Makarius Wenzel and Isabelle contributors. *Isar: Intelligible semi-automated reasoning*. Official project description. <https://isabelle.in.tum.de/Isar/>. Accessed September 5, 2026.