

Vantage

Mathematical Objects, Explicit Views, and Certified Computation

A representation-centered language over Lean,
with a formally checked computer-algebra core

A mathematician should be able to keep referring to the same object while changing how it is represented: as a finite sum, a generating function, a recurrence, a factored polynomial, or a locally defined analytic expression. Each change must say exactly what it preserves. Computation should produce such representations and their evidence, not silently replace the mathematical question.

Prepared for Vladimir Reshetnikov

5 September 2026, Pacific time

This is a technical design article, not a delivered language implementation. The repository review is static. Proposed syntax, interfaces, and metatheory are identified as specifications; the mathematical derivations are given explicitly. Companion exact-arithmetic checks are executable tests, not Lean formalizations. Repository revisions and the limits of the evidence are recorded in the article and the source manifest.

Abstract

The first brainstorming round converged on a Lean-hosted language of mathematical steps, scoped obligations, fixed statements, and replayable evidence. Its two synthesis reports identify the remaining difficulties particularly clearly: representation views do not yet have a satisfactory composition discipline; method contracts risk duplicating ordinary theorem types; definitions may create more friction than proof syntax; and a certificate that a theorem is true is not necessarily a certificate for the explanation printed beside it.

This article proposes VANTAGE, an object-and-view layer whose primary operation is *a certified change of representation for a specified observation*. The logical host remains Lean. The distinguishing design decisions are a separation of exact presentations, restricted observations, and one-way abstractions; support-sensitive composition without implicit strengthening; a typed protocol for exact, conditional, branch-complete, and finite-precision computational results; and a two-lane CAS architecture separating verified algorithms from untrusted certificate discovery. Normalization becomes query-directed rather than a single global normal form.

The design is grounded in ProveIt’s binomial inversion, autonomous derivative, and Bell generating-function developments, and in committed Leant source. A worked EGF interface shows how to avoid repeated factorial cancellation while retaining arbitrary commutative rational algebras. Further examples cover guarded rational simplification, root isolation and radical denesting, local differentiation, finite reindexing, and synthesis of proof connections. Concrete interfaces, failure cases, an implementation sequence, and responses to all twenty-two next-round questions in the two syntheses make the proposal testable. No reduction in authoring time, proof length, or maintenance cost is claimed as a measured result.

Contents

1	The proposal in one argument	5
1.1	What the first round has already settled	5
1.2	The three distinctions that carry the design	5
1.3	Why this is a proof assistant–CAS hybrid	5
1.4	What is new relative to the reports	6
2	Evidence from the repositories	6
2.1	Snapshots and scope of the review	6
2.2	Binomial inversion: the kernel identity, not the public corollary	6
2.3	Autonomous derivatives: locality is mathematical information	7
2.4	Bell generating functions: a definition-level opportunity	7
2.5	Leant: preserve the strongest boundaries already present	7
2.6	How much does the corpus audit establish?	8
3	An object-and-view model	8
3.1	Semantic anchors	8
3.2	Exact presentations	8
3.3	Observations and restricted information	9
3.4	Supports and hypotheses	9
3.5	No change to definitional equality	9
4	The language an author would write	10

4.1	A small stored language, not unrestricted prose	10
4.2	Four verbs with different effects	10
4.3	Calculations name the representation and the relation	11
4.4	Domain changes are lexical and explicit	11
4.5	Reasons: required, preferred, or unrestricted	12
4.6	References and incompleteness	12
5	View composition, reflection, and coherence	12
5.1	Exact presentations compose without choosing a global normal form	12
5.2	Conditional routes accumulate dependent supports	13
5.3	A three-view example: natural arithmetic through integers to rationals	13
5.4	Preservation is not reflection	14
5.5	Information-losing views need an observation budget	14
5.6	Germes are projections, not global function identities	14
5.7	An abstraction ledger is not an adequacy proof	14
6	Contracts and the typed intermediate representation	15
6.1	Logical contracts are ordinary theorems	15
6.2	Three kinds of method implementation	15
6.3	A common obligation record	15
6.4	IR nodes	16
6.5	What is checked, and what is merely reported	16
6.6	The modest assembly theorem	16
7	A verified computer-algebra architecture	17
7.1	Queries first, algorithms second	17
7.2	Lane A: verified algorithms and checked executions	17
7.3	Lane B: untrusted discovery with a verified checker	17
7.4	Reification is part of the mathematical boundary	18
7.5	Exact, conditional, covered, and approximate results	18
7.6	Branch coverage is substantive evidence	19
7.7	What belongs in the first verified algebra core	19
8	Worked algebra: cancellation, roots, and denesting	19
8.1	Polynomial certificates say exactly what they say	19
8.2	A rational-function identity is not a total-function identity	20
8.3	A radical identity with an explicit algebraic certificate	20
8.4	The same identity through algebraic-root presentations	21
8.5	Parameter-dependent radicals need branch-local statements	21
8.6	Certified approximations stay useful without becoming equalities	21
9	A complete EGF interface and the Bell example	22

9.1	The coefficient convention should be a named observation	22
9.2	Product and derivative rules	22
9.3	A representation-specific algebra, not an invisible instance switch	23
9.4	A coefficient compiler rather than a collection of ad hoc tactics	23
9.5	Represent the Touchard calculation as polynomial algebra	24
9.6	The shifted Bell identity	24
9.7	Spivey’s identity without the recurring factorial detour	25
9.8	What this saves, and what it does not	25
9.9	Finite computation and analytic interpretation remain separate	25
10	Local analysis and finite sums in the same model	26
10.1	Preserving the autonomous-derivative theorem	26
10.2	A failed local view should name the missing fact	26
10.3	Bounded sums provide scoped arithmetic facts	27
10.4	A finite-sum CAS can check telescoping certificates	27
10.5	Substantial analytic moves remain substantial	27
11	Integrating Leant without conflating its services	27
11.1	Two productive roles	27
11.2	An exact request boundary	28
11.3	Structural synthesis is useful even with opaque mathematics	28
11.4	Exact displayed-edit replay	29
11.5	Data synthesis needs a specified observation or behavior	29
11.6	Ranking should minimize burden, not only syntax	29
12	Definitions, packages, and collaboration	30
12.1	A definition package is an interface, not an automatic theory generator	30
12.2	The API should expose useful observations, not every representation detail	30
12.3	Instances and universe parameters stay explicit at the boundary	30
12.4	Versioning has two independent layers	30
12.5	Conflicting packages should not become a global rewrite soup	30
12.6	Migration from existing Lean	31
13	Publication, replay, and the reader’s evidence	31
13.1	One artifact, several projections	31
13.2	Search-free replay still performs computation	31
13.3	A practical trust boundary	32
13.4	Visibility should track mathematical decisions	32
13.5	Semantic notices are not missing hypotheses	32
14	Implementation sequence and performance	32
14.1	The smallest useful deliverable	32

14.2 Modules and ownership	33
14.3 Batching and sharing without weakening scope	33
14.4 Budgets describe the entire interaction	33
14.5 Implementation gates	34
15 Evaluation that separates language from library	34
15.1 Three baselines and a representation ablation	34
15.2 Tasks and controls	34
15.3 Leakage and infrastructure costs	35
15.4 Measures and decision criteria	35
16 Relation to existing work and limits	35
16.1 Relevant precedents, without claiming a new foundation	35
16.2 What is established in this article	36
16.3 The central risk and the useful fallback	36
17 Conclusion	36
A Answers to the syntheses' twenty-two questions	36
B A discriminating conformance suite	38
C Minimal surface grammar and result schema	39
D Companion checks and source provenance	41
D.1 Executed mathematical checks	41
D.2 Directly inspected files	41
References	42

1 The proposal in one argument

1.1 What the first round has already settled

Both synthesis reports recommend retaining Lean, expressing meaningful intermediate claims, giving omitted work exact scopes, fixing mathematical meaning before proof search, and retaining accepted evidence. They also require comparison with ordinary Lean equipped with the same helper libraries and synthesis services. These are the right starting constraints, not contributions that need to be invented again [1, 2].

The second round should answer a different question: *what mathematical abstraction should the shared infrastructure implement?* An obligation ledger records work, but does not by itself reduce the work. A convenient spelling of `have` makes a document easier to scan, but does not by itself explain why an expression can be differentiated, why a rational coefficient may be cancelled, or why a finite computation establishes an identity of infinite series.

My answer is to make *objects and certified representations* first-class. A mathematical argument typically retains an object while moving between presentations selected for different purposes. A polynomial is factored to locate roots, expanded to compare coefficients, and evaluated to establish a sign. An integer sequence is studied through a recurrence, a generating series, and a finite transform. A function is considered globally for a symmetry argument and locally for a derivative calculation. A useful language must preserve the object and make the change of viewpoint cheap, without confusing these different notions of preservation.

Central design decision. A computation or proof method does not merely return an expression or close a goal. It returns a representation of a fixed object, or evidence for a precisely named observation of that object, together with the hypotheses under which that result is valid.

1.2 The three distinctions that carry the design

An *exact presentation* denotes the same semantic object. A polynomial’s sparse coefficient list and its expression tree can be exact presentations of one polynomial. A *restricted observation* preserves only specified information: a Taylor jet preserves finitely many coefficients; a germ preserves local behavior; an interval enclosure preserves a containment fact. An *abstraction* supplies a directed logical implication and may lose information. Reducing an integer equation modulo a prime is generally an abstraction, not a reversible change of representation.

These distinctions are not cosmetic. If they are collapsed into one implicit “view” operation, the system can turn agreement of a finite prefix into equality of sequences, agreement at a point into agreement near that point, or an equality modulo a prime into an integer equality. Each of those is a different theorem.

The proposed language therefore keeps a stable semantic reference for the object, lets the author select a useful representation or observation, and asks the proof-producing infrastructure to construct the required bridges. This realizes the first round’s mathematical-step language without requiring every step to be a new bespoke tactic.

1.3 Why this is a proof assistant–CAS hybrid

The CAS portion operates on certified representations of exact mathematical objects. It can normalize polynomials, compute coefficients, propose factors, isolate algebraic roots, and calculate bounds. The proof-assistant portion handles quantifiers, arbitrary parameters, abstract structures, induction, local analytic hypotheses, and the assembly of those computations into larger arguments.

There is no special “CAS truth” admitted by the kernel. Algorithms either have proved correctness theorems and checked executions, or produce certificates for verified checkers. Their inputs and outputs use the same object, context, and observation protocol as ordinary theorem application. This is stronger than asking a commercial CAS for an answer and then hoping a general tactic can prove it. It is also more realistic than requiring every discovery heuristic to be verified before the environment can benefit from it.

1.4 What is new relative to the reports

The first round already contains theorem-backed views, contextual holes, semantic edit classes, and the separation of proof synthesis from data synthesis. VANTAGE adopts those ideas. Its proposed increment is to give them a more concrete mathematical and implementation boundary:

1. Exact presentations and information-losing observations have different interfaces, with explicit composition and reflection rules.
2. Computational results carry a precise specification, including domain conditions, branch coverage, and precision. Their evidence is tied to actual execution, not only to the correctness theorem of an algorithm.
3. A definition package supplies an observation interface, so recurring representation choices are resolved once rather than repaired in every proof.
4. Search is allowed to choose an implementation of an approved observation, but not its mathematical meaning. Leant can synthesize the connections between fixed interfaces, as well as individual proof terms.

These are design contributions to this campaign, not claims of historical priority. The relevant precedents are discussed in Section 16. The first implementation should be a Lean library and a small document layer. A separate surface language earns its place only if the controlled comparison supports it.

2 Evidence from the repositories

2.1 Snapshots and scope of the review

Both requested TeX reports were read closely, including their mutual corrections and next-round questions, at Brainstorm revision `494be9d5ab41a5ae219fc6f800e6aee3c93bce39`. This is important: the questions added to the second synthesis are included, not an earlier version of that report. The individual proposals are used through the two syntheses’ detailed crosswalks; this article does not claim an independent full rereading of all nine proposals.

The directly inspected ProveIt files are at public revision `24ce8bd743eaab64a91ce90725ea00f498d319d2`. The syntheses also discuss a newer local checkout, `7c4e3f10`; those observations are attributed to the syntheses rather than silently reassigned to the public snapshot. The inspected Leant revision is `3a40904be8a410d832d9ab6900b3d3e7b425eccb`. Some findings in the syntheses concern additional uncommitted working-tree changes. In particular, their discussion of a new behavioral **where** clause is not treated here as a fresh execution or as a feature verified in the committed source.

The review is static. Neither repository was built for this article. Bibliographic entries identify paths and immutable revisions. The companion manifest records which files and portions were inspected. This distinction matters more than a large count of files allegedly “reviewed.”

2.2 Binomial inversion: the kernel identity, not the public corollary

ProveIt’s `BinomialInversion.lean` deliberately works at two levels: arbitrary additive commutative groups acted on by integers, and commutative rings with integer coefficients cast into

them. The proof of kernel orthogonality passes through interval conversion, shifted indexing, a binomial product identity, natural-subtraction arithmetic, and cast normalization. The final inversion results then reuse a lower-triangular transform calculus [3].

The conclusion is not that the final theorem needs a more English-like spelling. It is that the kernel calculation would benefit from a stable finite-sum view. Inside the summand, the bound on the index should already be available to the arithmetic transport. The coefficient identity should be proved over integers and then transported forward, rather than rebuilding the calculation in every coefficient algebra. At the same time, replacing this development by a field-only CAS computation would lose one of its strongest properties: its original generality.

2.3 Autonomous derivatives: locality is mathematical information

`AutonomousODE.iteratedDeriv_eq_comp` is already a compact proof. In the successor case it converts the induction identity on an open set into an eventual equality near the current point, transports the derivative, applies the chain rule, and uses the recurrence. The hypotheses on the auxiliary functions are required only at values actually reached by the solution [4].

This is an excellent semantic test, not a good source of exaggerated compression claims. A frontend can present the argument as “differentiate the identity locally,” but it must retain the neighborhood information and the range-local hypotheses. Its benefit would be that the author handles a local function view, rather than repeatedly spelling out the representation of eventual equality. The source’s existing proof remains a control against overengineering.

2.4 Bell generating functions: a definition-level opportunity

The additional example examined here is `BellShiftEGF.lean`, together with the EGF interface in `BellGeneratingFunctions.lean`. It is particularly useful because the second synthesis identifies EGF-related files as promising places to investigate representation work [2, 5, 6].

The development represents an exponential generating function by an ordinary formal power series whose coefficient is

$$[X^n] \text{egfA}(a) = \rho(1/n!)a_n, \quad \rho : \mathbb{Q} \longrightarrow A.$$

Its abstract results work over an arbitrary commutative $\mathbb{Q}$ -algebra. To extract Spivey’s identity, the source specializes the generating-function identity to $\mathbb{Q}$, applies the ordinary coefficient function, moves a common $1/n!$ factor through a sum, proves that factor nonzero, cancels it, and reflects the resulting rational equality back to naturals [5].

The natural observation for that argument is not $[X^n]F$, but $n![X^n]F$. Giving that observation a proved interface removes the need to introduce and then cancel the factorial normalization at each use. This is a change of the definition’s public interface, not merely proof-text compression. Section 9 develops it completely and explains why it does not require making A a field or making ρ injective.

2.5 Leant: preserve the strongest boundaries already present

Leant’s committed `Verification.hs` explicitly says that its private `Verified` wrapper records acceptance by a supplied callback, not a behavioral proof, solver certificate, or kernel proof. The quota-aware scheduler groups spellings of a candidate, stops at accepted quotas, and does not copy acceptance authority from a merely equal-looking occurrence. These are useful boundaries for the new protocol [7].

The fragment translator retains structural connectives, supported constructor information, and typed application structure, while some dependent expressions remain opaque. Its comments and safety machinery distinguish qualified negative outcomes from unsuccessful bounded search [8].

The language should reuse that distinction without elevating a safety flag into a formal theorem about translation completeness.

Direct inspection of the committed suggestion path also confirms a narrower issue identified by both syntheses. A probe runs one tactic, extracts possible `Trythis` replacement text, and may classify the result by a reduction in visible goal count. The direct path does not independently replay that exact replacement before caching it. A separate completion receipt is therefore needed for the actual source edit, and current-goal progress must remain distinct from whole-proof completion [9]. This is an interface requirement for a sealed language, not a claim of a kernel soundness defect in an exploratory REPL.

2.6 How much does the corpus audit establish?

The second synthesis reports 90,637 classified tactic-like entries in 1,004 FabiusFunction files, with 29,303 entries in families it calls reconstructible. The first synthesis independently reproduced its full-subtree summary fields. Both revised discussions distinguish these lexical classifications from measured removable work [1, 2].

I use that audit to select examples, not to predict a percentage saving. A cast can be the mathematical substance of a theorem; a one-line solver call can hide many failed attempts; an assertion can hide a substantial proof. Likewise, no textual occurrence of `sorry` does not establish a successful build or an acceptable transitive axiom inventory. The experiment must measure actual obligations and author interventions, with semantic categories calibrated by inspection.

3 An object-and-view model

3.1 Semantic anchors

A semantic anchor is a resolved Lean object together with the environment and ordered context in which it is meaningful. At the specification level write

$$a = (\mathcal{E}, \Gamma, D, d), \quad \mathcal{E}; \Gamma \vdash d : D.$$

The word “anchor” describes a reference discipline, not a new mathematical type constructor in Lean. The same anchor may have several certified presentations. Names and pretty-printed strings are not identities. A stored anchor must bind universe parameters, type-class arguments, local instances, and the relevant declaration dependencies as well as the expression’s syntax.

The important restriction is that an algorithm cannot change D , d , or the meaning-bearing part of Γ just because a different problem is easier. It may propose a new anchor as an explicit definition or statement edit. Within a calculation it normally proposes a new representation of the existing anchor.

3.2 Exact presentations

For a fixed semantic type D , a presentation package has a representation type R , a validity predicate $\text{Valid} : R \rightarrow \text{Prop}$, and a meaning function

$$\text{meaning} : \{r : R \mid \text{Valid}(r)\} \longrightarrow D.$$

A presentation of $d : D$ consists of a valid r and an equality $\text{meaning}(r) = d$. The validity predicate can be trivial. When it is not, it contains genuine representation invariants: sorted exponents in a sparse polynomial, reduced denominators in rational data, or an isolating interval containing exactly one root of a nonzero polynomial.

It is often convenient to use a relation $V(d, r)$ instead of a meaning function. The functional form is preferable for the exact-presentation core because it makes the preserved object unambiguous.

Relational views remain available for objects whose representation is only specified up to an equivalence, but the relation and its transport laws must then be explicit.

A representation is not necessarily unique. Expanded and factored expressions may both denote the same polynomial. The language does not demand a unique canonical spelling for every mathematical object. It requires a certificate for the representation actually chosen.

3.3 Observations and restricted information

An observation is a typed operation on an object, or a proposition about it. Examples are the n -th normalized coefficient, evaluation at a point, root-count in an interval, an upper bound, or equality in a germ at a point. For a simple value-valued observation $q : D \rightarrow B$, a computation can return

$$b : B, \quad p : q(d) = b,$$

without constructing a new exact presentation of all of d .

Finite information requires this interface. A jet $J_N(F)$ of a formal series contains coefficients of degrees less than N . Its certificate says exactly that those coefficients agree with F . It is not a certificate that the truncated polynomial equals F . Similarly, an interval $[l, u]$ is a certified observation of a real value x when $l \leq x \leq u$, not a replacement for x by its midpoint.

One can describe a family of observations $\mathcal{O}$ by

$$d \equiv_{\mathcal{O}} e \iff \forall q \in \mathcal{O}, q(d) = q(e).$$

Transport through this relation is allowed only for registered operations proved to respect it. The relation is not inferred from the syntactic appearance of later uses. Adding a new observation can invalidate an old replacement receipt. For dependent observations, use a dependent transport theorem, not an untyped list of equal scalar outputs.

3.4 Supports and hypotheses

The support of a computational result records where, or under what assumptions, it is valid. Its logical content is an ordinary Lean proposition. Common supports include $x \in U$, $j \leq k \leq n$, $d \neq 0$, or a neighborhood equality. There is no separate unchecked “CAS assumption database.”

A conditional presentation can have the form

$$S \longrightarrow V(d, r).$$


It becomes available unconditionally in a local context only after S has been proved there. If S cannot be proved, the system may retain a conditional result, request a case split, or report that the selected operation is inapplicable. It may not append S to the theorem’s assumptions invisibly.

The support is not necessarily a nuisance. A nonzero denominator may be routine; a domination hypothesis for a limit interchange may be the key mathematical argument. The interface records their roles, but the theorem supplying the operation remains authoritative about what is required.

3.5 No change to definitional equality

VANTAGE does not make CAS simplification part of Lean’s definitional equality. Expressions related by polynomial normalization remain propositionally equal unless they were already definitionally equal in Lean. The elaborator inserts checked transports where needed.

This avoids making conversion depend on an expanding collection of algebra packages, nondeterministic search, or normalization budgets. It also prevents accidental global changes to instance resolution. The compiler may share and cache presentations aggressively, but the mathematical boundary is still an ordinary theorem about ordinary Lean expressions.



4 The language an author would write

4.1 A small stored language, not unrestricted prose

The initial frontend should reuse Lean expressions and binders, with a small set of document commands around them. Mathematical notation is a rendering of typed syntax. Free-form prose can explain an argument, but only formally marked claims are assertions certified by the system. An assistant may propose source edits; it does not interpret an unrestricted paragraph directly into trusted evidence.

The examples below are *proposed syntax*. Names such as `EGF.coefficients` and `rational.cancel` denote interfaces specified in this article, not existing commands claimed to compile. Ordinary Lean proof blocks remain valid leaves of the document and are checked under the same publication policy.

```

context (A : Type) [CommRing A] [Algebra Rat A]

object F : FormalSeries A := EGF A a

claim shifted : D F = EGF A (fun n => a (n + 1))
  by EGF.differentiate

observe c := ecoeff n F
  compute using EGF.coefficients

claim coefficient_value : c = a n
  from c.receipt

```

The `object` line introduces an ordinary definition with a stable semantic anchor. The `observe` line fixes the expression $\text{ecoeff}_n(F)$ before computation. Its result is not selected by searching for a conveniently provable coefficient theorem. The receipt certifies that the returned value is the value of that exact observation.

4.2 Four verbs with different effects

Claim introduces a proposition to prove. Any proof of that fixed proposition is logically acceptable, subject to its dependency and method policies. **Present** requests a new exact presentation of an existing object. **Observe** computes a specified observation without claiming an exact presentation of the whole object. **Select** chooses a new mathematical object with a specification; this is a data-producing operation and may change all later mathematics that refers to the selected object.

```

object roots : Set Real := {x | p.eval x = 0}

present roots as finite_algebraic_set
  compute using real_roots.complete

select r : Real
  satisfying p.eval r = 0 and 1 < r and r < 2
  using real_roots.unique_in_interval

```

The first computation owes both soundness and completeness of the returned root set. The second owes a witness in the stated interval. If uniqueness is part of the selection method, it also owes a uniqueness theorem. Merely producing one root is insufficient for the first command. Merely producing any root is insufficient for the second.

An existential proof used inside another proof is a different operation: `obtainrwithhfromexistence`. It performs scoped elimination in `Prop`. It is not automatically an executable definition of a root. Construction of a dependent pair in `Type`, propositional existence, and a noncomputable choice operator retain their ordinary Lean distinctions.

4.3 Calculations name the representation and the relation

```

claim cancellation (x : Real) (hx : x != 1) :
  (x*x - 1) / (x - 1) = x + 1
calculate over Real with support hx:
  (x*x - 1) / (x - 1)
  = ((x - 1) * (x + 1)) / (x - 1) by polynomial
  = x + 1 by rational.cancel

```

The first equality is a ring identity followed by congruence; the second requires the nonzero denominator. The local support supplies that premise. A command without `hx` can discover the conditional result, but cannot publish this unconditional equality for an arbitrary x .

The same expression can be explored without asserting the missing premise:

```

observe q := (x*x - 1) / (x - 1)
  compute using rational.cancel

-- Result, not an unconditional rewrite:
--   under x != 1: q = x + 1
--   unresolved support: x != 1

```

A publication can retain the conditional theorem as such. What it cannot do is print $q = x + 1$ while keeping the premise only in a hidden warning.

4.4 Domain changes are lexical and explicit

A block such as `nearx`, `overRat`, or `coefficientsbelowN` changes the working observation or representation, not the original object. The elaborated view shows the map used to enter the block and the theorem used to leave it. There is no ambient convention that all functions are smooth or that all denominators are nonzero.

```

claim derivative_step : HasDerivAt f v x
  near x:
    present f as g using local_identity
    derive g has_derivative v using chain_rule
    conclude by local_derivative_transport

```

Here `local_identity` must establish equality in a neighborhood, and the last line names a theorem respecting that local relation. A pointwise equality at x has the wrong type of information. The frontend can hide the proof that an open set is a neighborhood, but cannot hide the absence of such a proof.

4.5 Reasons: required, preferred, or unrestricted

A named transformation in a `calculate` block is method-constrained by default. If the author says “reindex by this map,” the method receipt must contain that map, its coverage, and its summand correspondence. A different proof of the equality is not a successful implementation of that explanation.

For an ordinary claim, `bysearch` means unrestricted reconstruction within the stated policy. A `prefer` clause influences candidate order but does not assert that the preferred method was used. The published rendering names the actual accepted method. Changing a required reason is a visible plan edit, not a silent tactic fallback.

This gives explanation a tractable meaning: the displayed method is the instantiated proof constructor actually recorded. It does not pretend to prove that the method was logically indispensable or that the proof reproduces a human’s internal reasoning process.

4.6 References and incompleteness

Authored claims receive stable identifiers. A printer may render a reference as “the preceding identity” only when that phrase resolves to the stored target. Inserting a new identity does not silently retarget an accepted reference. Copying a block creates fresh identities and checks the new scope.

Drafts may contain open claims, conditional computations, unsupported methods, and proposed data choices. They are useful mathematical working documents. Publication checks every formal assertion, including unused ones. A proof of the final theorem cannot compensate for an unproved intermediate sentence. Informal commentary is visibly distinguished from certified assertions rather than being made to look equally checked.

5 View composition, reflection, and coherence

5.1 Exact presentations compose without choosing a global normal form

Suppose R_i, R_j, R_k present the same semantic type D , with meaning maps m_i, m_j, m_k . A conversion $f : R_i \rightarrow R_j$, on valid inputs, has a theorem

$$m_j(f(r)) = m_i(r).$$

A second conversion $g : R_j \rightarrow R_k$ has $m_k(g(s)) = m_j(s)$. Ordinary equality transitivity yields

$$m_k(g(f(r))) = m_i(r).$$

This is the basic composition rule. The intermediate representation is concrete and its invariants must be checked before the next conversion uses it.

Proposition 5.1 (Anchor-relative coherence). *Let two accepted routes from an anchored object $d : D$ produce presentations $r : R_i$ and $s : R_j$, with proofs $m_i(r) = d$ and $m_j(s) = d$. Then every observation $q : D \rightarrow B$ has*

$$q(m_i(r)) = q(m_j(s)).$$

Proof. The two meaning equalities give $m_i(r) = m_j(s)$. Applying q to that equality gives the conclusion. No equality between the representation types or their inhabitants is required. $\square$

This is a useful, limited answer to the syntheses' coherence question. Two factorization algorithms may return factors in different orders. If both presentations denote the same polynomial, an observation of the polynomial is unaffected. An observation of the order of the factor list is not an observation of the polynomial unless a separate convention has made that order part of the specified object. The latter is a data choice, and must not be silently changed.

5.2 Conditional routes accumulate dependent supports

The real problem is often conditional conversion. Write a route as

$$r_0 \xrightarrow[G_1]{f_1} r_1 \xrightarrow[G_2]{f_2} r_2 \xrightarrow[G_3]{f_3} r_3.$$

Here G_2 can mention r_1 , and G_3 can mention both the value and invariants produced at earlier stages. The correct record is an ordered telescope, not a set of pretty-printed predicates:

$$h_1 : G_1(r_0), \quad h_2 : G_2(r_0, h_1, r_1), \quad h_3 : G_3(r_0, h_1, r_1, h_2, r_2).$$

The instantiated route is checked as one dependent program. If the supports are independent propositions, a conjunction is an adequate presentation. The implementation must not assume that special case universally.

Guard generation does not grant a guard. It creates a proof obligation in the appropriate context. Failed alternatives are rolled back before another route is tried, including assignments to shared metavariables. After a route is accepted, the concrete intermediate representations, support proofs, and method versions are retained.

5.3 A three-view example: natural arithmetic through integers to rationals

Let $a, b, d \in \mathbb{N}$, assume $b \leq a$, $d > 0$, and $d \mid (a - b)$. Consider

$$u = (a \dot{-} b) \operatorname{div} d \in \mathbb{N},$$

where both operations are the natural operations. A route may first transport $a \dot{-} b$ to its integer difference, next embed in $\mathbb{Q}$, and finally represent exact division by rational division. It owes, in order, the bound needed for subtraction, the divisibility fact needed for exact division, and the nonzero denominator needed for the rational manipulation.

A safe implementation can reason through a witness k with $a - b = dk$. It proves that the natural quotient is k , and that the rational quotient of the casts is also the cast of k . The equality then follows through that common witness. It does not identify floor division with field division. For $a = 6, b = 0, d = 4$, the cast of the natural quotient is 1, while the rational quotient is $3/2$. For $a = 2, b = 3$, the natural difference is 0, not -1 .

A direct natural-to-rational route and the route through integers may coexist. If both have the same anchored input, exact target observation, and discharged supports, Proposition 5.1 gives observational agreement. The recorded route still matters for explanation, performance, and maintenance. Adding a cheaper route must not change a source-selected representation or silently discard one route's necessary support.

5.4 Preservation is not reflection

A homomorphism $h : A \rightarrow B$ preserves an equality: $a = a'$ implies $h(a) = h(a')$. To conclude $a = a'$ from the equality in B , injectivity or an appropriate reflection theorem is needed. A ring normalization in a quotient ring cannot establish an arbitrary equality in the original ring.

The adapter interface consequently distinguishes

$$\text{preserve} : P \rightarrow Q, \quad \text{reflect} : Q \rightarrow P.$$

Either direction may exist without the other. A $\mathbb{Q}$ -algebra map is useful for forward transport of rational identities even when no injectivity hypothesis is available. The natural-number embedding into $\mathbb{Q}$, in contrast, supplies the reflection used at the end of a natural-number coefficient argument.

This distinction must be part of method typing, not only documentation. The compiler may search for a chain of directed bridges, but it may not reverse an arrow because the displayed formulas look similar.

5.5 Information-losing views need an observation budget

For a formal series F , let

$$J_N(F) = ([X^0]F, \dots, [X^{N-1}]F).$$

Equality of $J_N(F)$ and $J_N(G)$ permits comparison of coefficients below N , not arbitrary coefficients. Multiplication preserves this precision: $J_N(FG)$ depends only on $J_N(F)$ and $J_N(G)$. Differentiation consumes one coefficient of precision: $J_{N-1}(F')$ is determined by $J_N(F)$ when $N \geq 1$.

Substitution has its own contract. The simple truncated substitution rule is valid when the inner series has zero constant coefficient: terms of outer degree at least N cannot contribute below degree N . More general admissible substitutions, for example involving nilpotent constant terms, require a different precision theorem. Merely knowing that the host's substitution is defined does not establish the simple zero-constant precision rule.

Thus a view records not just its result type but its *available observations*. A demand for coefficient N cannot be fulfilled from a jet of length N . The system requests more precision or another representation. It does not turn finite evidence into an infinite identity.

5.6 Germs are projections, not global function identities

Passing from a function $f : \mathbb{R} \rightarrow \mathbb{R}$ to its germ at x forgets behavior away from x . An eventual-equality certificate identifies two germs, not two global functions. The derivative-transfer theorem is an observation theorem respecting that germ relation. The route from a global function to a germ is therefore not entered into the exact-presentation graph for global functions.

The distinction also prevents a subtle error with domain restrictions. A formula valid only on U can give a germ formula at x when U is a neighborhood of x . A membership proof $x \in U$, by itself, does not do so. The transition needs a neighborhood theorem, which openness may supply.

5.7 An abstraction ledger is not an adequacy proof

Both syntheses rightly emphasize that negative results need more care than positive candidates. One refinement is essential: a list of abstracted nodes is useful diagnostic metadata, but neither its emptiness nor a provider's claim of complete search is a proof of end-to-end adequacy.

Suppose a provider searches an abstract proposition A for a target P . A checked proof of A transfers positively through a checked theorem $A \rightarrow P$. A checked proof of $\neg A$ transfers negatively through a checked theorem $P \rightarrow A$. Those are different arrows. Each must account for the exact context and foundational policy.

There is a further distinction between *non-derivability in a proof calculus* and a proof of a negation. Complete intuitionistic propositional proof search can fail to derive an arbitrary instance of $P \vee \neg P$, but that failure cannot be imported as $\neg(P \vee \neg P)$. Indeed intuitionistic reasoning already proves

$$\neg\neg(P \vee \neg P) : \quad h \mapsto h(\text{Or.inr}(p \mapsto h(\text{Or.inl } p))).$$

A syntactic search-completeness certificate is therefore not automatically an object-level refutation. Leant’s fragment-qualified outcomes should remain qualified metadata unless an actual proof at the original negated Lean goal is constructed. This proposal strengthens the provider boundary without alleging that Leant currently makes this particular invalid inference.

6 Contracts and the typed intermediate representation

6.1 Logical contracts are ordinary theorems

The second synthesis’s proposal to use annotated theorem types is adopted, with one adjustment to its durability story. Binder names are better than positional flags, but they are not stable identities across library refactorings. A project-owned adapter should expose stable semantic roles and map them to the current theorem through a typechecked wrapper [2, 1].

For example, a finite reindexing adapter exposes **source**, **target**, **map**, **covers**, **injective**, and **summand**. A new upstream lemma name or reordered binder changes the wrapper, not every proof document. The wrapper itself is ordinary Lean and its proof is charged to the shared library. The logical type is never duplicated in an unchecked contract language.

Roles control authoring and presentation. They do not decide logical relevance. A bijection proof is not automatically routine because its type lives in `Prop`; it may be the point of the argument. The method policy can attempt it automatically, but records both the proposition and what was used.

6.2 Three kinds of method implementation

Theorem adapters instantiate one theorem or a small composite theorem. Their unsupplied binders become obligations. **Certified computations** reify a supported input, run an algorithm or checker, and return evidence for a specified relation. **Structural plans** use recursors, case splits, and reductions with scoped subdocuments. Their motives and coverage are explicit.

All three return typed evidence, but they should not be forced into the same implementation API. An attribute on a theorem is enough for the first kind, not for the entire mechanics of a normalizer or induction elaborator. Conversely, introducing a new symbolic method object for every ordinary lemma would add unnecessary maintenance.

6.3 A common obligation record

The authoritative worker owns the elaborated expressions. An obligation is conceptually

$$O = (id, origin, \mathcal{E}, \Gamma, T, roles, policy, dependencies).$$

Here Γ is an ordered dependent telescope, T is the fixed expected type, and the origin includes the source revision, anchor identities, and worker generation. There are separate controls for available local facts, library constants, and allowed transitive axioms.

The provider receives a closed request over that telescope or opaque handles owned by the worker. Pretty-printed formulas may accompany the request for retrieval or model prompting, but do not replace its identity. Restricting the local context must retain the dependencies required to type the remaining variables. It is not equivalent to deleting unwanted printed hypothesis lines.

6.4 IR nodes

A small typed document IR needs the following node families. The table describes interfaces, not a proposed change to Lean’s kernel.

Node	Meaning	Required retained information
Definition	Introduce an anchored term.	Type, value, instances, exported observations.
Assertion	Prove a fixed proposition.	Context, exact target, evidence, source identity.
Presentation	Denote the same object in a new representation.	Representation invariant, meaning equality, support.
Observation	Compute a property or value of an anchor.	Query, precision/domain, output specification.
Selection	Construct a new specified object.	Exact witness, behavioral specification, choice policy.
Method	Apply a required theorem adapter or computation.	Instantiated contract, actual inputs and sub-obligations.
Case/induction	Introduce scoped subdocuments.	Cover or recursor, motive, local contexts, all branches.
Lean leaf	Admit an ordinary Lean proof body.	Verbatim source plus its checked boundary and dependencies.

A document is not simply a DAG while elaboration metavariables remain unresolved. There may be shared data and instance constraints across several nodes. Such a cluster is elaborated transactionally; its closed interfaces are frozen before independent provider work begins. Evidence dependencies are then acyclic after recursors are treated as checked term constructors, not circular proof promises.

6.5 What is checked, and what is merely reported

A receipt can report a discovered candidate, an accepted local proof, a checked computational specification, method conformance, or successful independent replay. These statuses are not interchangeable. In particular, a successful callback is not promoted into a universal specification theorem.

Final document acceptance requires a proof of the exact theorem, evidence for every asserted intermediate claim, typechecked definitions and selections, accepted computational specifications, and the declared dependency policy. Method-constrained nodes additionally require their advertised constructor. Human explanatory value and fidelity of informal commentary remain separate questions. The renderer must never display a single undifferentiated green state that conceals these distinctions.

6.6 The modest assembly theorem

After the interfaces are closed, each node is an ordinary Lean term constructor applied to checked predecessors. Dependent substitution gives a term at the fixed final target; the final kernel check is retained deliberately. This is the same relative soundness argument accepted by both syntheses, and it does not need another elaborate metatheory chapter [1, 2].

The interesting obligations lie in preserving that argument’s premises: exact input identity, coherent supports, no scoped witness escaping, correct reification, checked computation, complete

branch coverage, and faithful source-to-claim mapping. Sections 5 and 7 specify those boundaries more closely. They are paper specifications here, not a formal proof of an implemented compiler.

7 A verified computer-algebra architecture

7.1 Queries first, algorithms second

A typed query fixes an input type I , output family $O(i)$, precondition $\text{Pre}(i)$, and relation

$$\text{Spec}(i, o) : \text{Prop}.$$

Examples include polynomial equality, a factorization with product equal to the input, the complete set of real roots, a particular coefficient, or an enclosure. The user-visible command chooses the specification. Providers choose how to produce a value satisfying it.

“Simplify” is too vague as the fundamental query. A factorization query can specify equality of the product, but irreducibility is an additional property. A normal-form query can specify a canonical ordering, but an equality query need not require canonical output. A root query can request all roots or one selected root. Different specifications require different certificates even if their outputs look superficially similar.

7.2 Lane A: verified algorithms and checked executions

Suppose a pure algorithm has been defined and verified:

$$\text{alg} : I \rightarrow O, \quad \text{alg_correct} : \forall i, \text{Pre}(i) \rightarrow \text{Spec}(i, \text{alg}(i)).$$

This does not, by itself, certify an arbitrary value o returned by a native executable called `alg`. One also needs a checked connection between that execution and the logical function: for example

$$e : \text{alg}(i) = o,$$

or a verified evaluator, refinement theorem plus justified evaluation, or a computation trace whose verified checker establishes the same relation.

This execution boundary is easy to overlook when saying “based on formally verified algorithms.” Proving the source algorithm correct is one obligation; showing that the value actually used is its result is another. VANTAGE’s strict profile requires both. Small calculations can use kernel reduction. Larger calculations can use compact certificates or a verified trace checker. A more permissive native-computation profile can exist, but its additional trust is explicit rather than being advertised as the strict profile.

Lean’s current validation documentation makes this distinction operationally relevant: native computation can introduce computation-specific axioms, and since Lean 4.29 the historical single `Lean.trustCompiler` name is no longer a sufficient audit target. The policy must inspect the actual transitive axiom inventory, not blacklist one remembered constant [11].

7.3 Lane B: untrusted discovery with a verified checker

For difficult discovery tasks, define a certificate type $C(i, o)$, a total checker, and a soundness theorem:

$$\begin{aligned} \text{check} &: \prod_i \prod_o C(i, o) \rightarrow \text{Bool}, \\ \text{check_sound} &: \forall i, o, c, \quad \text{check}(i, o, c) = \text{true} \rightarrow \text{Spec}(i, o). \end{aligned}$$

An external CAS, Leant, a language model, or a specialized search program may propose (o, c) . Its authority ends there. The exact output is accepted only through a checked computation of the checker and its soundness theorem. Malformed, oversized, or unsupported data is rejected before expensive checking.

The checker need not be complete: rejecting a valid but unsupported certificate is an operational limitation, not a mathematical refutation. Its own cost must also be bounded by input and certificate limits. Certificate checking can be expensive, especially for expanded polynomials, and should not be called cheap without measurement.

Both lanes can implement one query. A verified factorization algorithm and an untrusted factor generator can return the same factorization specification. The author sees the same mathematical result but can inspect different execution and provenance receipts. This permits incremental development without confusing a verified checker with a fully verified discovery engine.

7.4 Reification is part of the mathematical boundary

An algebra backend usually does not compute on arbitrary Lean expressions. It computes on a reflected syntax of polynomials, rationals, indices, or series operations. The reifier must supply the connection to the original expression:

$$\text{denote}(e, \nu) = t,$$

where e is reflected syntax, ν is the exact typed interpretation of its variables, and t is the anchored Lean term. The environment records the actual coefficient ring and its operations, not just the printed name “ring.”

For example, treating an uninterpreted real expression as an atom is safe for a positive ring identity when its exact interpretation is retained. Treating noncommuting matrices as commuting atoms is not. Erasing natural subtraction into ring subtraction is not. Approximating an algebraic number by a floating point constant is not an exact reification.

After computation, the result is decoded through a checked interpretation theorem and assembled into the original target. A certificate for the wrong reified expression is useless even when its internal arithmetic is impeccable.

7.5 Exact, conditional, covered, and approximate results

The mathematical result family is richer than success/failure:

```
Exact(output, proof_of_specification)
Conditional(guard_telescope, output, proof_under_guards)
Covered(branches, proof_of_coverage, branch_specifications)
Enclosure(bounds, error_or_containment_proof)
Unsupported(feature, source_node)
Exhausted(resource_observations)
```

These are the checked-side concepts. An untrusted provider sends candidate payloads, not authority-bearing `proof_of_specification` fields that the consumer simply believes.

A conditional result proves $G \rightarrow \text{Spec}(i, o)$. A covered result contains conditions $G_1, \dots, G_m$, a proof that the current assumptions imply their disjunction, and a result under each condition. If branches overlap and the query specifies a unique value, consistency follows from the specification’s uniqueness theorem; otherwise a deterministic selection policy must be part of the output. A piecewise result must use only data available in its branch.

An enclosure establishes a relation such as $l \leq x \leq u$, possibly with a specified rational error. It is not an equality to a decimal approximation. The display may round for readability, but the

certificate stores exact bounds. A request for a strict sign may be discharged by an enclosure not containing zero; an enclosure containing zero is inconclusive, not a proof of zero.

7.6 Branch coverage is substantive evidence

Consider the equation $ax = b$ over a field. The usual expression b/a solves it when $a \neq 0$. A complete solution-set computation must also handle $a = 0, b = 0$, where every x is a solution, and $a = 0, b \neq 0$, where none is. The equality $a = 0 \vee a \neq 0$ and the corresponding split on b are part of the proof of coverage. A “generic case” answer is valid only as a conditional answer until those exceptional branches are accounted for.

This generalizes to discriminants, parameter-dependent root multiplicities, vanishing leading coefficients, and singular matrices. The language should show these as mathematical branches, not obscure CAS warnings. Which branch an algorithm finds first is not a convention for interpreting the theorem.

7.7 What belongs in the first verified algebra core

Start with exact integers and rationals, sparse polynomial operations, polynomial evaluation and differentiation, finite coefficient convolutions, and small certificate checkers for algebraic identities. These support the EGF and reindexing examples without a general symbolic integration engine. The next layer can add polynomial division with remainder, Euclidean algorithms, root-count certificates, and rational interval arithmetic. Symbolic integration, multivariate real quantifier elimination, and general transcendental identities are later providers with their own precise domains. The language protocol does not promise a universal algorithm for them. It makes unsupported cases and conditional results useful rather than misleading.

Existing Lean and Mathlib automation should be reused where its evidence fits. The current Mathlib `Polyrith.lean` documentation, for example, says that the external Sage-based `polyrith` path is no longer supported after the server shutdown and points to `grobner` as an alternative for polynomial goals. A roadmap must not depend on the old hosted service being available. The certificate interface itself should be backend-independent [12].

8 Worked algebra: cancellation, roots, and denesting

8.1 Polynomial certificates say exactly what they say

A useful algebraic certificate for a consequence of equations $f_1 = \dots = f_m = 0$ is an identity

$$p = \sum_{i=1}^m q_i f_i.$$

A verified polynomial normalizer can check the identity. Evaluating it in a commutative ring proves that $p = 0$ whenever the f_i vanish. The discovery engine may use a Gröbner basis, but the checker need not repeat that discovery.

This certificate establishes ideal membership. It does not by itself provide a complete decision procedure for every semantically valid consequence over every ring. An implication valid over the reals may also use order, radicals, or nonzero factors. A more general certificate

$$d^k p = \sum_i q_i f_i$$

proves $p = 0$ only with suitable cancellation evidence for d^k . In a field, $d \neq 0$ suffices. In an arbitrary ring, a nonzero element can be a zero divisor, so regularity or unit evidence is needed. The query must name the ambient algebraic structure.

The factorization specification likewise has levels. Checking that $p = c \prod_i f_i^{e_i}$ certifies a factorization. Certifying that each f_i is irreducible requires additional evidence in a specified coefficient domain. The display must not upgrade a product identity into a complete irreducible factorization merely because an external routine calls its result “factor.”

8.2 A rational-function identity is not a total-function identity

In the fraction field $\mathbb{Q}(X)$, the identity

$$\frac{X^2 - 1}{X - 1} = X + 1$$

is unconditional because $X - 1$ is a nonzero polynomial. When interpreted as a formula in a real variable with total real division, the corresponding equality is conditional on $x \neq 1$. At $x = 1$, the left-hand side is 0, whereas the right-hand side is 2.

The fraction-field computation is still useful. It produces a presentation of a rational function, plus an evaluation theorem on points where the denominator does not vanish. The language keeps these as two distinct steps. It does not pretend that equality in a fraction field is equality of every totalized pointwise evaluation.

For differentiation at a point $x_0 \neq 1$, continuity of $x \mapsto x - 1$ supplies a neighborhood on which the denominator stays nonzero. Only then may the rational-function simplification be used as a local function identity and differentiated. This is a three-way interaction among algebra, support, and analysis that a useful hybrid should handle without hiding the argument.

8.3 A radical identity with an explicit algebraic certificate

Consider

$$\sqrt{3 + 2\sqrt{2}} = 1 + \sqrt{2}.$$

Put $u = \sqrt{2}$ and $v = \sqrt{3 + 2u}$. The real-root specifications provide $u^2 = 2$, $u \geq 0$, $v^2 = 3 + 2u$, and $v \geq 0$. The algebraic part is the polynomial identity

$$(v - u - 1)(v + u + 1) = (v^2 - 3 - 2u) - (u^2 - 2). \quad (1)$$

The right side is zero, and $v + u + 1 > 0$. Cancellation therefore gives $v - u - 1 = 0$, which is the desired result.

This is a particularly transparent CAS certificate: a polynomial identity plus a sign fact. It shows why algebraic equations alone do not identify the intended root. Replacing v by its negative still satisfies the squared equation but violates the nonnegative-root specification. The denominator in the cancellation certificate is not an incidental formatting detail; it selects the correct consequence of the polynomial equations.

```
claim denesting : sqrt (3 + 2 * sqrt 2) = 1 + sqrt 2
  let u := sqrt 2
  let v := sqrt (3 + 2*u)
  have hu : u*u = 2 and 0 <= u by real_sqrt
  have hv : v*v = 3 + 2*u and 0 <= v by real_sqrt
  have balance : (v-u-1)*(v+u+1) = 0
    by polynomial_certificate using hu hv
  conclude by cancellation using balance and positivity
```

A direct square-and-nonnegativity method would be equally reasonable. The system should not use a root isolator for this example merely to demonstrate its most complicated subsystem. The certificate vocabulary makes both proof routes available and their actual costs comparable.

8.4 The same identity through algebraic-root presentations

For a more general denesting or algebraic-number calculation, use a polynomial and an isolating interval. Here both sides are roots of

$$P(X) = X^4 - 6X^2 + 1$$

in $(2, 3)$. A Sturm sequence is

$$P(X), \quad 4X^3 - 12X, \quad 3X^2 - 1, \quad \frac{32}{3}X, \quad 1. \quad (2)$$

Each term after the first two is the negative remainder of its two predecessors. At 2, the signs are $(-, +, +, +, +)$, with one variation. At 3, they are all positive, with no variation. The endpoints are not roots. A root-count theorem therefore gives exactly one root in $(2, 3)$.

For $u = \sqrt{2}$, the inequalities $1 < u < 2$ place $1 + u$ in $(2, 3)$, and substitution using $u^2 = 2$ proves $P(1 + u) = 0$. For $v = \sqrt{3 + 2u}$, we have $4 < 3 + 2u < 9$, so $2 < v < 3$; its square equation gives $P(v) = 4u^2 - 8 = 0$. Uniqueness in the interval proves the identity.

An exact real-algebraic presentation can thus store a nonzero integer polynomial, rational endpoints, a root-count certificate, and the unique-root meaning. Polynomial irreducibility is not required merely to identify a unique root in an interval. Zero polynomials, endpoint roots, repeated roots, and the distinction between distinct-root counts and multiplicities need explicit treatment in the checker. They are not safely handled by assuming a generic square-free case.

A full real-root-set result needs a global completeness certificate as well as individual isolating intervals. For example, a verified root bound and counts on a disjoint interval decomposition can show that every real root was included. A list of individually valid roots establishes only soundness of the list, not completeness. Multiplicity annotations add another specification layer.

8.5 Parameter-dependent radicals need branch-local statements

For real x , $\sqrt{x^2} = |x|$, not unconditionally x . A branch-complete presentation is x on $x \geq 0$ and $-x$ on $x < 0$, together with coverage. For complex roots, even this order-based split is unavailable; a principal-branch convention and its branch-cut conditions belong to the semantic object.

Consequently, a denesting provider should first state whether it is returning an identity in an algebraic extension, a real-valued branch identity on a domain, or a principal complex identity. These are not interchangeable outputs of the same untyped simplifier. The initial implementation should support exact real algebraic constants before attempting general parameterized complex radicals.

8.6 Certified approximations stay useful without becoming equalities

Numerical search can help isolate a root, propose an interval, or find a likely factor. Its output is then checked by exact arithmetic. For a transcendental function, a bracket $[a, b]$ plus continuity and opposite endpoint signs proves existence of a root; a positive lower derivative bound on the interval can prove uniqueness and a residual-to-error estimate.

For instance, if $f(r) = 0$, $r, x \in [a, b]$, and $f' \geq m > 0$ throughout the interval, the mean value theorem gives

$$|x - r| \leq \frac{|f(x)|}{m}.$$

A rigorous interval evaluation of $f(x)$ then gives a certified enclosure for r . The numerical algorithm suggests x ; the theorem and exact bounds supply the claim. The proof requires the interval and derivative hypotheses, not just a small floating-point residual.

9 A complete EGF interface and the Bell example

9.1 The coefficient convention should be a named observation

Let A be a commutative $\mathbb{Q}$ -algebra and let $\rho : \mathbb{Q} \rightarrow A$ be its structure map. Define

$$u_n = \rho(1/n!), \quad v_n = \rho(n!).$$

Because $n! \neq 0$ in $\mathbb{Q}$, the rational identity $(1/n!)n! = 1$ transports forward to

$$u_n v_n = v_n u_n = 1_A. \quad (3)$$

No injectivity of ρ is needed for this equation. In particular, these arguments do not need to exclude a trivial coefficient ring. They use an explicit inverse identity, not cancellation from a claim that a mapped scalar is nonzero.

For a sequence $a : \mathbb{N} \rightarrow A$, set

$$E(a) = \sum_{n \geq 0} u_n a_n X^n, \quad \text{ecoeff}_n(F) = v_n [X^n]F.$$

These are the semantic definition and observation that the EGF presentation package should expose.

Proposition 9.1 (EGF coordinates over arbitrary commutative rational algebras). *For all a and $F \in A[[X]]$,*

$$\text{ecoeff}_n(E(a)) = a_n, \quad E(n \mapsto \text{ecoeff}_n(F)) = F.$$

Thus the EGF coordinate map is a bijection between sequences and formal series.

Proof. The first identity is $v_n u_n a_n = a_n$, by (3). For the second, the coefficient of X^n on the left is $u_n v_n [X^n]F = [X^n]F$. Formal-series extensionality gives equality. $\square$

This is a representation theorem with a genuine inverse, not an assumption that the coefficient map into A reflects equalities from $\mathbb{Q}$. The proofs of rational scalar identities take place in $\mathbb{Q}$, and their images are used in A .

9.2 Product and derivative rules

Define binomial convolution by

$$(a \star b)_n = \sum_{k=0}^n \rho\left(\binom{n}{k}\right) a_k b_{n-k}.$$

Then

$$E(a)E(b) = E(a \star b), \quad \text{ecoeff}_n(FG) = \sum_{k=0}^n \rho\left(\binom{n}{k}\right) \text{ecoeff}_k(F) \text{ecoeff}_{n-k}(G). \quad (4)$$

To check the scalar factor in the second equation, use the rational identity

$$n! \frac{1}{k!} \frac{1}{(n-k)!} = \binom{n}{k} \quad (0 \leq k \leq n)$$

and transport it through ρ . The finite-sum binder supplies the index bound; it does not need to be rediscovered at each multiplication node.

For the formal derivative D , the coefficient formula $[X^n]DF = (n+1)[X^{n+1}]F$ gives

$$\text{ecoeff}_n(DF) = \text{ecoeff}_{n+1}(F), \quad DE(a) = E(n \mapsto a_{n+1}). \quad (5)$$

This follows from $n!(n+1) = (n+1)!$, again in the scalar map. Repeated application shifts by any m , with no analytic convergence question.

The source already proves the corresponding EGF product and derivative laws. The proposed change is to make normalized coefficients an explicit, reusable observation interface and to let the coefficient compiler work directly in it, not to present those identities as new mathematics [6].

9.3 A representation-specific algebra, not an invisible instance switch

The sequence type $\mathbb{N} \rightarrow A$ also has its familiar pointwise operations. Under EGF coordinates, however, series multiplication becomes $\star$, not pointwise multiplication. The language must not silently replace the ring instance on raw sequences when a proof would become easier.

Use named operations or a distinct wrapper type, such as a structure whose data field is a sequence and whose multiplication is explicitly binomial convolution. The presentation package identifies that algebra. A statement written before entering the package does not acquire a different multiplication afterward. The resolved-statement view shows which operation is being used.

A possible package declaration is:

```
presentation EGFCoordinates (A) on FormalSeries A
  representation := EGFSquence A
  meaning a := egfA A a.coefficients
  observation normalized_coefficient n := factorial n * coeff n
  laws := coefficient_inverse, series_inverse
  operations :=
    add      : pointwise_add
    multiply : binomial_convolution
    derive   : shift
```

Every listed law is an ordinary theorem. The package declaration organizes those laws; it does not make a multiplication convention true by declaration.

9.4 A coefficient compiler rather than a collection of ad hoc tactics

A first coefficient compiler can recognize constants, sums, scalar multiples, products, formal derivatives, and $\exp(aX)$. It recursively compiles the requested normalized coefficient into finite algebra, accompanied by applications of (4)–(5) and the corresponding base rules. Unsupported subexpressions remain typed atoms or explicit obligations.

For example,

$$\text{ecoeff}_n(C(c) \exp(aX) E(b)) = c \sum_{k=0}^n \rho \left(\binom{n}{k} \right) b_k a^{n-k}.$$

Here $C(c)$ is the constant series. The author chooses the coefficient observation; the compiler supplies routine scalar transport and finite convolution. It is not asked to discover Spivey's identity by arbitrary theorem search.

A compiled rule is a theorem schema for all n , not a finite table of tested coefficients. For numerical coefficient queries, a precision-indexed jet is often more efficient. The same observation protocol distinguishes these two modes: a symbolic universal theorem and a computation at one fixed index.

9.5 Represent the Touchard calculation as polynomial algebra

Write $U = \exp(X)$ and let

$$T_m(Y) = \sum_{j=0}^m S(m, j) Y^j,$$

where $S(m, j)$ are Stirling numbers of the second kind. Their recurrence gives

$$T_0(Y) = 1, \quad T_{m+1}(Y) = Y(T'_m(Y) + T_m(Y)). \quad (6)$$

For $1 \leq j \leq m+1$, the coefficient on the right is $jS(m, j) + S(m, j-1) = S(m+1, j)$; the zero coefficient and the upper endpoint are checked separately using the usual zero-extension convention. These endpoints are part of the polynomial representation's support law.

The formal derivative obeys $DU = U$. Hence

$$D(T_m(U)) = UT'_m(U), \quad D(T_m(U)) + UT_m(U) = T_{m+1}(U).$$

The long derivative-and-shift calculation in `derivative_touchardExp` can therefore be approached through an ordinary polynomial representation followed by a proved evaluation/derivative bridge [5].

This is an example of query-directed form selection. The finite exponential sum is useful for extracting coefficients; the polynomial in U is useful for proving the derivative recurrence. There is no reason to insist that one be the canonical form for both jobs. The exact presentation receipts connect both to the same formal series.

No injectivity of polynomial evaluation at U is needed here: the polynomial identity is transported forward to a series identity. Inferring an arbitrary polynomial identity backward from equal evaluations would require additional evidence and is not part of this method.

9.6 The shifted Bell identity

Let

$$B(X) = \sum_{n \geq 0} B_n \frac{X^n}{n!}$$

in EGF notation, interpreted through ρ in A . The Bell recurrence establishes $DB = UB$. Induction gives

$$D^m B = T_m(U)B. \quad (7)$$

The base case is $T_0 = 1$. If the identity holds at m , then

$$\begin{aligned} D^{m+1} B &= D(T_m(U))B + T_m(U)DB \\ &= (UT'_m(U) + UT_m(U))B = T_{m+1}(U)B. \end{aligned}$$

This is exactly the mathematical structure preserved by the source proof. The proposed frontend should expose this induction and recurrence, not hide them inside a command named after the final theorem.

```
claim shifted_bell (m : Nat) : D^[m] B = T m (exp X) * B
induct m:
  zero:
    conclude using T_zero
  successor m using ih:
    differentiate ih formally
    use B_derivative and Touchard.polynomial_recurrence
    conclude by commutative_algebra
```

The notation $D^\sim[m]$ here means iteration of the formal derivative, not pointwise exponentiation. This interpretation is fixed in the expression layer. It has no connection to analytically differentiating a convergent power series unless a separate analytic bridge is invoked.

9.7 Spivey’s identity without the recurring factorial detour

Apply ecoeff_n to (7). The left side is B_{n+m} by (5). On the right,

$$T_m(U) = \sum_{j=0}^m S(m, j) \exp(jX), \quad \text{ecoeff}_r(\exp(jX)) = j^r.$$

The coefficient compiler therefore obtains

$$B_{n+m} = \sum_{j=0}^m \sum_{k=0}^n S(m, j) \binom{n}{k} B_k j^{n-k}. \quad (8)$$

As an identity in A , the integers and naturals are interpreted through their canonical scalar maps. To conclude the stated natural-number identity, instantiate $A = \mathbb{Q}$, normalize casts, and use the injectivity of the natural embedding into $\mathbb{Q}$. This last reflection is a separate bridge. It is not a property assumed for every $\mathbb{Q}$ -algebra.

```
claim spivey (m n : Nat) :
  bell (m+n) = sum j in [0..m], sum k in [0..n],
    stirling2 m j * choose n k * bell k * j^(n-k)
work in Rat by faithful_nat_embedding:
  observe normalized_coefficient n of shifted_bell m
  present T m (exp X) as its finite exponential sum
  conclude by EGF.coefficients
```

The omitted work has a bounded, compositional specification: evaluation of a polynomial as a sum, a finite coefficient convolution, rational scalar identities, and cast reflection. The theorem statement and the finite bounds stay visible. The source’s use of a commutative rational algebra has not been replaced by a stronger global field hypothesis.

9.8 What this saves, and what it does not

The repeated factorial manipulation is moved into an ordinary library interface whose laws are proved once. The polynomial presentation similarly centralizes finite-support and evaluation plumbing. The induction and the choice to use a generating function remain authored mathematics.

This is a plausible reduction in representation work, not a measured authoring improvement. Ordinary Lean given the same ecoeff and polynomial interfaces may become just as concise. The experiment must measure that baseline. If it does, the correct result is a better library plus an optional readable document view, not a claim that new syntax caused the improvement.

9.9 Finite computation and analytic interpretation remain separate

Checking (8) for finitely many (m, n) is a useful regression test. It does not establish the universal identity. Computing a jet of B can answer specific coefficient queries, but not (7) for all orders. And a formal identity in $A[[X]]$ is not automatically an equality of real functions at every x .

A future analytic-series package needs convergence and evaluation theorems with explicit domains. It may consume a formal identity as one input, but the convergence hypotheses are additional obligations. The language must not make the word “series” select whichever interpretation lets the proof close.

10 Local analysis and finite sums in the same model

10.1 Preserving the autonomous-derivative theorem

Let $U \subseteq \mathbb{R}$ be open, and suppose $W, \varphi : \mathbb{R} \rightarrow \mathbb{R}$ satisfy

$$W'(x) = \varphi(W(x)) \quad (x \in U)$$

in the sense of an actual derivative-existence statement. Let G_n, H_n be families such that, for $x \in U$,

$$G_0(W(x)) = W(x), \quad G'_n(W(x)) = H_n(W(x)), \quad G_{n+1}(W(x)) = H_n(W(x))\varphi(W(x)).$$

Only derivative existence at the displayed range points is needed. The conclusion is $W^{(n)}(x) = G_n(W(x))$ for every n and $x \in U$, matching the relevant generality of the ProveIt theorem [4].

The induction invariant is the identity on all of U , not just at the current point. At $x \in U$, openness turns that identity into equality of germs. Differentiating the representative $G_n \circ W$ gives

$$(G_n \circ W)'(x) = H_n(W(x))\varphi(W(x)) = G_{n+1}(W(x)).$$

Derivative transport through the germ equality gives the successor case.

```
claim all_orders : forall n x, x in U -> D^[n] W x = G n (W x)
induct n retaining the identity on U:
  zero:
    conclude using G_zero
  successor n using ih:
    fix x in U
    near x using U_is_open:
      present D^[n] W as G n o W using ih
      differentiate using G_derivative_at_image and W_equation
      use G_successor
    conclude by local_derivative_transport
```

The author-visible decisions are induction, retaining the identity on U , and using a local composition. The generated work is neighborhood construction, function composition bookkeeping, and exact theorem application. The system does not introduce global smoothness of G_n , nor replace an arbitrary W by a canonical solution it happens to know how to differentiate.

10.2 A failed local view should name the missing fact

Replace the invariant on U by equality at one point. The local presentation fails with a diagnostic of the form:

```
Cannot enter the requested germ presentation at x.
Available: f x = g x.
Required: f and g agree on a neighborhood of x.
The derivative method has not been attempted.
```

This is much more useful than a failed general search for the final derivative. For $f(t) = 0$, $g(t) = t$, and $x = 0$, the available point equality is true while the derivative conclusion is false. The diagnostic explains the exact semantic gap. A proposed repair that adds a neighborhood hypothesis is a statement change unless it can be proved from the original assumptions.

10.3 Bounded sums provide scoped arithmetic facts

For integers n, j represented as naturals with $j \leq n$, put $m = n - j$. The bijection $i \mapsto j + i$ between $0 \leq i \leq m$ and $j \leq k \leq n$ gives

$$\begin{aligned} \sum_{k=j}^n (-1)^{n-k} \binom{n}{k} \binom{k}{j} &= \binom{n}{j} \sum_{i=0}^m (-1)^{m-i} \binom{m}{i} \\ &= \binom{n}{j} [m = 0] = [n = j]. \end{aligned}$$

The complementary case $n < j$ is an empty interval. The language should not pretend it is covered by the same natural-subtraction normalization.

Inside the transformed summand, $0 \leq i \leq m$ is part of the local telescope. It supplies the index facts needed by binomial identities and casts. Reindexing owes a bijection or multiplicity-aware alternative, and equality of corresponding summands. A half-open replacement $0 \leq i < m$ fails coverage at the upper endpoint. This is a failure of the advertised method even if a broad solver can prove the final theorem by an unrelated route [3].

10.4 A finite-sum CAS can check telescoping certificates

A useful additional finite method accepts a proposed antidifference G and checks, on the exact summation range,

$$f(k) = G(k+1) - G(k).$$

It then derives

$$\sum_{k=a}^b f(k) = G(b+1) - G(a)$$

for the selected finite interval convention. A discovery engine may guess G or compute it symbolically; the checker verifies the difference identity and endpoint conditions. If G is rational, every denominator used at the interior and boundary points must be covered by the support.

This method demonstrates a genuine CAS–proof-assistant division of labor. The expensive discovery of an antidifference can be untrusted, while its use in an inductive or combinatorial proof remains a small verified operation. It grants no right to telescope an infinite series without the additional limiting theorem.

10.5 Substantial analytic moves remain substantial

Dominated convergence, differentiation under an integral, and analytic continuation should initially be theorem adapters with explicit hypotheses. The language can organize an integrable dominating function and the relevant local uniformity facts, but cannot classify their existence as routine by default.

For example, the pointwise limit of $f_n = n\mathbf{1}_{(0,1/n)}$ is zero on $[0,1]$, while each integral is one. Pointwise convergence alone therefore cannot justify exchanging limit and integral. A good hybrid can compute those integrals and verify the counterexample; it cannot obtain the invalid interchange by a CAS rewrite rule. This is exactly where expressive supports matter more than terse syntax.

11 Integrating Leant without conflating its services

11.1 Two productive roles

Leant has two relevant jobs: proposing proof-state edits and synthesizing typed terms. The new environment should not collapse them into a generic “prove” endpoint. A tactic edit changes an

elaboration state; a synthesized term is checked at a fixed expected type. Both can help, but they carry different replay requirements [9, 8].

For VANTAGE, the first integration target is term assembly across fixed interfaces. Suppose the local context contains a neighborhood membership proof, an identity on that neighborhood, and a derivative certificate for a representative. The adapter library exposes ordinary theorems connecting these to the desired derivative claim. Leant can synthesize the composition without being asked to invent the neighborhood or change the theorem’s domain.

The second target is bounded route synthesis. Candidate terms assemble presentation converters and observation transport theorems, with their support proofs. The route endpoints and direction are fixed. A successful route is rechecked as a term in the authoritative Lean worker. Failure remains failure to find a route, not evidence that the mathematical claim is false.

11.2 An exact request boundary

The Lean-side adapter owns `Exprs` and context declarations. A request to an external provider contains a stable request token, immutable origin, allowed provider inventory, a structural search representation, opaque atom handles, and resource limits. The worker keeps the authoritative association between handles and expressions. Rendered formulas are auxiliary data.

Request	
<code>origin</code>	: <code>environment + source_generation + obligation_id</code>
<code>context</code>	: ordered typed telescope, owned by Lean worker
<code>expected</code>	: exact type handle
<code>effect</code>	: <code>evidence representation observation selection</code>
<code>inventory</code>	: explicitly allowed theorem/adaptor handles
<code>limits</code>	: total time, memory, steps, candidate and payload caps
<code>abstraction</code>	: diagnostic map plus any checked transfer contracts
Candidate	
<code>request</code>	: original request token
<code>construction</code>	: typed assembly recipe or candidate source
<code>proposed_edit</code>	: optional exact source replacement
<code>observations</code>	: search and resource metadata, without proof authority

The final candidate is elaborated or reconstructed against the original expected type. No implicit declaration is added to make an unknown identifier meaningful. Universe and instance arguments that determine the target’s meaning have already been resolved or generalized explicitly. Candidates cannot reinterpret an opaque atom by supplying a new pretty-printed string with the same spelling.

11.3 Structural synthesis is useful even with opaque mathematics

A synthesis engine need not understand real analysis to combine $H_1 \rightarrow H_2$, $H_2 \rightarrow H_3$, and a term of H_1 . It can treat the analytic propositions as opaque typed atoms. This is precisely where structural synthesis complements a CAS. The CAS checks a polynomial or coefficient relation; Leant finds a short typed route from that relation and the contextual premises to the next claim.

Dependent arguments remain a boundary. A route that chooses an index map or a new witness may affect the types of later obligations. The worker must either instantiate those arguments first or validate the whole dependent candidate as one transaction. A product-of-independent-holes encoding cannot be assumed for all mathematical methods. Unsupported dependent structure should remain visible rather than be erased to obtain a false completeness claim.

There is no reason to rewrite the Haskell engines before measuring this use. A Lean-native adapter owns meaning; the existing engines remain replaceable candidate generators. The comparison with `exact?` and ordinary theorem application should distinguish single-lemma lookup from genuine

compositions. If nearly all accepted cases are the former, a dedicated adapter may not justify its cost.

11.4 Exact displayed-edit replay

For tactic suggestions, the accepted transaction is the *exact edit proposed to the author*, not the tactic whose diagnostic message happened to contain it. The protocol is:

1. Probe a candidate in a disposable worker cloned from an immutable origin.
2. Extract the exact replacement source, preserving scope and layout.
3. Replay that replacement from the same origin in a clean proof state.
4. Record its actual residual obligations or its completed proof term.
5. Accept the edit only if the live document still has the same origin.

A response with fewer visible goals means progress or closure of a local goal, not automatically completion of the entire theorem. The checked result records remaining goals, unresolved metavariables, admissions, root-term status, and publication policy separately. Missing status fields cannot be interpreted as success. A late response after undo may be saved as a stale proposal, but cannot mutate the current proof.

This addresses the exact boundary seen in the committed `suggestTactic` path without discarding its useful behavior: looking past the first nonclosing candidate, trying short composed continuations, caching by origin, and preserving the user’s script after backend failure [9].

11.5 Data synthesis needs a specified observation or behavior

For a proof of a fixed proposition, different accepted terms are normally interchangeable as evidence. For a selected datum, type inhabitation alone is not a behavioral contract. Two projections both inhabit $A \rightarrow A \rightarrow A$; a constant failure function can inhabit the type of an option combinator. A candidate must meet the property that downstream mathematics actually uses.

Finite tests are useful for exploration but certify only their finite propositions. A universal behavior claim remains a separate proof obligation. The syntheses’ discussion of Leant’s experimental behavioral work motivates this separation, but is not a report that this article executed that work [1, 2].

Not every witness inside a proposition needs user interaction. A proof of an existential proposition may construct a witness automatically while keeping the public proposition fixed. The stronger stability rule applies when the document names a particular witness as part of its mathematical argument or exports it as observable data. The accepted witness and its specification are then retained; substituting a different value is not merely replacing a proof term.

11.6 Ranking should minimize burden, not only syntax

For representation and route candidates, record a cost vector rather than a single unexplained score: unresolved mathematical choices, new supports, number of conversions, expected checking cost, certificate size, and adherence to the requested explanation. Deterministic project policies can order this vector. The system should not prefer a ten-character answer requiring a new field hypothesis over a longer proof in the stated ring.

Proof-only holes can usually stop at the first accepted candidate within policy. For data selections with genuinely different allowed values, present meaningful alternatives or require a canonical specification. The published artifact needs the selected object and its evidence. Rejected alternatives belong in an optional exploration log, not the logical dependency closure of the theorem.

12 Definitions, packages, and collaboration

12.1 A definition package is an interface, not an automatic theory generator

The proposed definition package contains an ordinary Lean definition, explicit operation conventions, exported observations, representation invariants, and proved bridge theorems. It also names simplification orientations and the methods permitted to construct alternative presentations. It does not generate arbitrary mathematics from a natural-language definition.

The first package should be the EGF coordinate interface of Section 9. Its scope is small enough to implement and rich enough to exercise nontrivial instances, finite supports, formal derivatives, and faithful versus nonfaithful scalar transport. A second package can handle polynomials and real-algebraic constants. A general controlled-language definition compiler should wait for evidence that these interfaces cannot be expressed conveniently with ordinary Lean declarations and templates.

12.2 The API should expose useful observations, not every representation detail

For an EGF, normalized coefficients, formal derivative, and binomial convolution are natural observations and operations. For a bounded function, point evaluation and its bound may be the useful interface, while the internal subtype projections are not. For an algebraic root, evaluation of a polynomial and interval membership are useful, whereas the order in which a solver found its isolating interval is not.

Hiding representation details is safe only through proved interface theorems. A user can explicitly open an implementation view when those details matter. That creates new dependencies. For example, reasoning about the precise list of factors means later reordering that list can no longer be treated as invisible. The distinction is analogous to an abstract data type, but the interface laws are mathematical theorems and are available to proof reconstruction.

12.3 Instances and universe parameters stay explicit at the boundary

The semantic seal includes the resolved instances used by the object and its observations. A new ring instance or different coercion route can change the meaning even when the printed formula remains unchanged. A proved proposition used as a local instance, such as a property of an already fixed structure, is different from selecting a different operation-bearing structure.

Unresolved universe constraints are also not left for a remote search engine to settle opportunistically. A polymorphic declaration generalizes the intended universe parameters. A monomorphic query fixes them. The worker may reconstruct implicit proof arguments, but not change a value's type universe or algebraic interpretation because one candidate elaborates more conveniently.

12.4 Versioning has two independent layers

A logical adapter version identifies the theorem and its typed role interface. An operational policy version identifies search bounds, normalization orientation, ranking, rendering, and permissible automatic supports. One can change while the other remains unchanged.

A project-owned wrapper gives stable role names across upstream refactoring, but its actual type is still checked. Reusing the same role identifier with a different type does not make old evidence compatible. The old environment remains pinned; a migration constructs new evidence and compares the resulting semantic anchors and theorem statements.

12.5 Conflicting packages should not become a global rewrite soup

Imports make packages available; they do not automatically activate every normalizer and coercion policy in them. A source block selects an active profile lexically or by a qualified method name.

Two conflicting normalization conventions require an explicit selection. Combining rewrite sets indiscriminately can change termination, output form, and explanatory behavior even when each rewrite is true.

Restrictions on permitted dependencies and axioms can be intersected when profiles are composed, provided the remaining context is still well typed. Capabilities may be combined only after the active interpretation and method priorities have been resolved. “Last import wins” is not an acceptable semantic rule. The resolved profile is part of the receipt, so a collaborator sees which policy was used without reconstructing an accident of import order.

12.6 Migration from existing Lean

A lifting tool can recover named claims from `have` blocks, calculations, and local definitions, attach retained evidence, and render them as a document. It should preserve opaque Lean leaves when it cannot recognize a method safely. It cannot generally recover an author’s intended explanation from a proof term. A recognized method annotation must itself be checked against the extracted proof boundary, or remain a proposed annotation.

The initial round-trip guarantee should cover a documented subset: anchored claims, explicit definitions, calculations, and the specified method nodes. Arbitrary metaprogram source is preserved verbatim rather than normalized into invented prose. This gives an adoption path that does not require rewriting the entire ProveIt corpus before obtaining a useful view of it.

13 Publication, replay, and the reader’s evidence

13.1 One artifact, several projections

A published VANTAGE proof bundle is specified to contain the source document, an expanded Lean companion, the anchored statement and definition manifest, accepted step evidence, computation certificates, package versions, the dependency and axiom audit, and a bidirectional source map. The human argument, obligation view, and exact evidence view are projections of those records, not independently maintained documents.

A hash identifies content and supports cache lookup. It is not a proof. On a cache hit, the expected type, context identity, semantic anchor, and certificate schema are checked before evidence is reused. A proof from a sibling case must not be imported because its printed goal matches the current one.

13.2 Search-free replay still performs computation

Replay does not rerun theorem retrieval, CAS discovery, synthesis ranking, or model calls. It checks retained terms and certificates against the pinned contracts and statements. Evaluating a verified certificate checker is still computation; “search-free” does not mean zero checking cost.

Some large results may be best stored as compact certificates rather than expanded proof terms. Others are cheap to retain as terms with shared dependencies. The protocol permits both while recording which execution boundary is used. Replaying an arbitrary saved tactic is a different mode because that tactic may perform fresh search or depend on changed simplifier state.

The source remains the authoring representation. In the original pinned environment, the accepted term or checked certificate is the evidence authority. After an upgrade, the old evidence remains a historical artifact; it is not silently treated as valid in a new environment. Migration produces a new receipt, compares statement and definition meanings, and reports changes to objects or plans. There is no promise that serialized Lean internals are portable across all versions.

13.3 A practical trust boundary

Lean already separates elaboration from its small logical checker; VANTAGE uses that architecture instead of inventing a new foundation [10]. Untrusted producers run outside the acceptance worker. The strict publication profile checks exact targets, all formal assertions, computation evidence, and a transitive axiom allowlist. For externally supplied executable proof code, process isolation protects the checking environment as well as logical state.

Current Lean documentation describes independent rechecking and a comparator workflow that checks against a separately supplied statement. These are relevant precedents, not features this article claims to have implemented or run. A local prototype can start with ordinary Lean checking and an explicit status report, while preserving an interface that can support the stronger deployment boundary [11].

13.4 Visibility should track mathematical decisions

The default reader view shows the theorem, the main claims, induction motives, selected objects, representation changes, and substantial support assumptions. Routine proofs of already visible bounds can be collapsed. The derivative example should show why local equality is available; the denesting example should show which root is selected; the EGF example should show that coefficients are normalized and formal rather than analytic.

Clicking a step expands its exact instantiated contract, obligations, their sources, and the accepted evidence. The system can report that a true intermediate claim was not used. It cannot silently drop an unproved assertion to make a publication succeed. A method-constrained rendering states what operation was checked, not a stronger claim that no alternative proof exists.

13.5 Semantic notices are not missing hypotheses

A statement using natural subtraction or total real division is not automatically wrong. The renderer should disclose its interpretation. A cancellation command, on the other hand, may owe a nonzero premise. These are different statuses: a semantic notice invites meaning review; a contract obligation requires evidence.

A dismissal of a notice can be scoped to a statement or to an explicit project convention. Its key includes the relevant expression and convention version, so a changed operation does not inherit an unrelated dismissal. New notices should be sparse and evaluated for false alarms. A linter that warns about every division teaches the author to stop reading rather than improving fidelity.

14 Implementation sequence and performance

14.1 The smallest useful deliverable

The first deliverable is not a standalone parser with dozens of mathematical verbs. It is a Lean package that implements anchored queries, ordinary theorem adapters, a source-linked receipt, and two representation packages: normalized EGF coefficients and exact polynomial expressions. Ordinary Lean and a minimal step frontend consume the same interface.

The end-to-end acceptance case is one coefficient proof: author a generating identity, apply the normalized coefficient observation, reconstruct the finite algebra, retain evidence, and replay with discovery disabled. Its negative neighbor changes the requested coefficient beyond the available jet precision. This is a small enough slice to reveal whether the object-and-observation model is actually useful.

14.2 Modules and ownership

Module	Responsibility	Must not decide
Core anchors	Resolved terms, contexts, identities, closed query boundaries.	Which theorem the author meant.
Document IR	Claims, source maps, scope, method and selection nodes.	That an unchecked assertion is harmless.
Presentation registry	Meaning maps, observations, support and transport laws.	That a one-way map is reversible.
Method adapters	Stable role interface to ordinary Lean theorems.	That every proposition-valued input is routine.
Algebra core	Reflected syntax, exact operations, verified checking.	That numeric resemblance is equality.
Provider scheduler	Budgets, candidate streams, cancellation, stale origins.	Proof or specification acceptance.
Leant bridge	Structural translation and origin-preserving candidate exchange.	A new interpretation for an opaque atom.
Evidence publisher	Exact-target validation, certificates, dependencies, replay.	Cross-version compatibility without rechecking.
Renderer/editor	Meaning disclosures, views, editable proposed patches.	Semantic identity from pretty-printing.

14.3 Batching and sharing without weakening scope

Reifying the same coefficient expression for every tiny obligation would be wasteful. Cache reflected subexpressions and interpretation proofs inside one fixed environment and context. Use DAG sharing for polynomial and series expressions, and retain common scalar identities as ordinary lemmas.

For a coefficient calculation, batch finite scalar identities over the same ring and index telescope. Do not batch across incompatible branch assumptions. If obligations share unresolved data choices, resolve or transactionally check that cluster first. Scoped caches carry the origin, active instances, profile, and precision. A certificate at precision N can answer smaller coefficient queries through a proved restriction rule, never larger ones by wishful reuse.

Normalization should be demand driven. Expanding $(X + 1)^{100000}$ is a poor way to answer its degree or leading-coefficient query. The observation method can prove those facts directly from a compact representation. A factorized polynomial may remain factorized until a particular coefficient requires an expansion. The proposed economy is semantic, not a claim that all expensive algebra becomes cheap.

14.4 Budgets describe the entire interaction

Per-tactic heartbeats are not an end-to-end deadline. The request budget covers candidate enumeration, reification, subprocess execution, communication, certificate parsing, checking, and exact displayed-edit replay. Limits include memory, output size, number of candidate attempts, and branch count as well as elapsed time.

The scheduler returns partial *information* without partial *proof acceptance*: a checked conditional result, an accepted intermediate presentation, or a located unsupported subexpression can be useful even when the full query exhausts its budget. Cancellation cannot turn a missing final

receipt into a successful proof. Worker generations prevent stale results from being committed after undo or a changed definition.

14.5 Implementation gates

Gate A: one exact path. The normalized-coefficient proof works through ordinary Lean and through the minimal document frontend, using the same library. Every output observation has a checked specification. Replay performs no discovery.

Gate B: semantic neighbors. Add finite jets, guarded rational evaluation, and the local derivative adapter. Reject precision overreach, a lost denominator condition, and pointwise-only equality with diagnostics naming the missing fact.

Gate C: a CAS certificate. Check the denesting polynomial certificate and a univariate root-count example. Mutated signs, omitted roots, invalid interval endpoints, and missing branch coverage must fail at their exact boundary.

Gate D: synthesis. Connect Leant for closed structural assembly and bounded view routes. Run exact-source replay and stale-origin tests before advertising verified tactic edits. Compare the closure and latency profile with existing Lean alternatives.

Gate E: authoring evidence. Run the frontend and definition-package comparison, including repair after library changes. Expand the grammar only when a failed task demonstrates a need that cannot be met by a shared library interface or a better document rendering.

These gates are deliverables, not a calendar estimate. They deliberately leave large symbolic integration and general natural-language interpretation out of the first implementation.

15 Evaluation that separates language from library

15.1 Three baselines and a representation ablation

Use three primary conditions: idiomatic existing Lean; refactored Lean with the same adapters, observation packages, CAS checkers, and Leant access; and the new frontend over that same machinery. The difference between the last two measures the surface and interaction layer. The difference between the first two measures shared mathematical and automation infrastructure.

Add a representation ablation: perform the EGF tasks both with the original ordinary-coefficient interface and with normalized coefficients, keeping the frontend fixed. This tests the definition-package thesis independently of the syntax thesis. An improved ordinary Lean proof is a successful outcome, not a result to hide because it weakens a language marketing claim.

15.2 Tasks and controls

Use the inspected binomial and derivative examples as development tasks, and Bell-related coefficient and recurrence tasks as a richer representation test. Hold out entire theorem families or modules before tuning adapters. Include at least one already-short corollary as a control: the new frontend should be allowed to preserve its short Lean form.

A reading study asks participants to identify the central argument, locate an essential assumption, and predict the effect of a mutation. An authoring study records time, edits, unsuccessful attempts, explicit mathematical choices, and new helper lemmas. A maintenance study changes an endpoint, a coefficient convention, an instance, a theorem name, and a package version separately. Frontend order should be counterbalanced, and Lean and domain experience recorded.

15.3 Leakage and infrastructure costs

Remove the target theorem and its downstream consequences from the permitted retrieval environment. A cached proof, alias, example, or provider wrapping the target is still leakage. Split by theorem or module before harvesting subgoals, not after near-identical goals have entered the training or tuning corpus. Record the exact environment and dependency exclusions for every task.

Charge for all new mathematics, wrappers, checker proofs, parser work, and maintenance. For n tasks, report

$$C(n) = L + \sum_{i=1}^n (A_i + R_i),$$

where L is shared infrastructure cost, A_i is authoring effort, and R_i is repair effort under the defined mutations. Do not amortize L over hypothetical future users. Report the components and observed reuse.

15.4 Measures and decision criteria

Source length, generated term size, and bridge counts are useful descriptions, not universal utility measures. Record median and tail interaction latency, checking time, certificate size, support obligations needing author input, reader accuracy, and repair effort. Publish failed tasks and false-positive notices as well as successful demonstrations.

Exact statement preservation, scope safety, precision discipline, and the required negative-case behavior are hard correctness gates. Usability thresholds should be chosen after a pilot provides actual variance, not invented in this article. If normalized coefficients help but new syntax does not, ship the coefficient library. If the document view helps readers but slows authors, offer it as a rendering. If the method registry costs more to maintain than its saved work, reduce its scope rather than conceal the cost in generated files.

16 Relation to existing work and limits

16.1 Relevant precedents, without claiming a new foundation

The immediate intellectual source is the nine-proposal campaign as evaluated by its two syntheses. The claim-and-obligation layer, statement locks, method conformance, and fair shared-library baseline are adopted explicitly, not presented as new discoveries [1, 2].

CoqEAL is a relevant precedent for connecting mathematical representations with more efficient computational ones through refinement. Its project separates abstract algebraic material, computation-oriented data structures, and a refinement framework. VANTAGE proposes a language-facing organization around that general kind of representation discipline, not the invention of verified representation changes [13].

Li, Passmore, and Paulson’s work on univariate polynomial problems in Isabelle/HOL separates external certificate discovery from formally checked real-algebraic reasoning. The associated Sturm–Tarski formalization is a concrete precedent for the root-count lane. These are Isabelle developments, not evidence that a matching Lean implementation is already present. Their execution and code-generation trust boundaries must be assessed in their own setting; citing them does not supply a strict Lean execution receipt [14, 15].

The design’s distinctive emphasis is the shared source-level protocol: an algebraic certificate, a local analytic identity, and a normalized coefficient computation all report what observation they establish of which fixed object, on what support. It is an integration and interface proposal over an existing logical foundation.

16.2 What is established in this article

The algebraic and EGF derivations are supplied as mathematical arguments. The Sturm-chain arithmetic and small finite identities are independently exercised by the companion exact-arithmetic program. The repository observations are static source readings at recorded revisions. The design provides concrete contracts and tests, not an implemented compiler or a Lean-verified metatheory.

In particular, no claims are made here that the proposed syntax compiles, that the new CAS checkers have been formalized in Lean, or that authors become faster. No ProveIt or Leant build or live backend acceptance suite was run. The companion program is a regression aid and certificate illustration, not a replacement for the verified checker that the implementation would need.

16.3 The central risk and the useful fallback

The main risk is not that Lean lacks a place to check the evidence. It is that creating and maintaining observation interfaces costs more than the authoring friction they remove, or that the interfaces hide the very distinctions readers need. The scoped examples and controlled evaluation are designed to test that risk directly.

The useful fallback is substantial: normalized coefficient interfaces, stable adapters, exact suggestion replay, and better source-linked evidence can all improve ordinary Lean without a new language becoming the primary input form. The proposal is therefore not contingent on inventing a more English-looking syntax. Its technical hypothesis is that mathematical objects deserve explicit, certified, query-oriented representations.

17 Conclusion

A concise proof language should not ask a mathematician to forget whether an expression is a polynomial, a formal series, a local function, or an algebraic root. It should make it easy to use each object in the representation appropriate to the argument, while preserving exactly what that representation justifies.

VANTAGE keeps the first round’s conservative logical foundation and makes the representation boundary its central language abstraction. Exact presentations preserve objects; restricted observations carry precision and domain; directed abstractions require the correct transport arrow. Verified algorithms and certificate checkers answer the same typed queries, with execution evidence attached to the actual outputs. Leant assembles typed connections and proposes edits without controlling mathematical meaning.

The Bell example supplies a concrete first experiment: expose normalized EGF coefficients, reuse polynomial representations for recurrences, and compare both frontends with the same proofs and services. The derivative and root examples supply the nearby failures that prevent this economy from becoming semantic ambiguity. That is the proposed second-round contribution: fewer representation choices repeated by hand, with more precise contracts for every choice that still matters.

A Answers to the syntheses’ twenty-two questions

The identifiers A1–A10 refer to the questions in *Mathematical Ideas, Checked Evidence*; B1–B12 refer to those in *Nine Blueprints*. The answers below state a provisional decision and an observation that could change it, rather than repeating the questions as another open-ended roadmap.

ID	Question area	Decision and discriminating artifact
A1	Smallest useful step	An ordinary theorem adapter, except for computation and structural plans. Implement one reindexing adapter and one normalized-coefficient query; add richer state only where these require it.
A2	Does a reason constrain?	Named calculation moves are required methods; ordinary claims can use unrestricted search or preferences. Compare a valid alternate proof with a failed advertised method.
A3	Automatic choices	Proof evidence and hidden exact representations may change within policy; named witnesses, chosen structures, plans, and statements are retained. Mutate a root sign and an instance separately.
A4	Verified suggestion	The exact displayed edit is independently replayed from its origin. Local progress, whole-proof completion, and publication are separate receipts. Test stale undo, invalid suggestion text, and missing status.
A5	Infrastructure cost	Begin with measured EGF, polynomial, finite-sum, and locality interfaces. Give the same packages to Lean; record helper development and held-out reuse, not only shorter examples.
A6	View coherence	Exact routes share a semantic anchor; limited observations have separate transports. Run the guarded natural–integer–rational route and add a competing direct route. Compare supports and preserved observations.
A7	Statement review	Render domains, quantifier dependencies, selected operations, branches, and precision. Notices are not obligations. Measure detection of nearby changed statements and false alarms.
A8	Retained artifact	Store author source plus checked terms/certificates. The latter certify the pinned environment; migration creates new evidence. Compare a rename with a changed definition, not just rebuild success.
A9	Definition authoring	Implement an EGF coordinate package with named operations and inverse laws. Compare ordinary coefficients with normalized observations through both frontends, charging the package cost.
A10	Published visibility	Show mathematical choices and substantial supports; collapse their routine derivations. Reader tasks must locate the local-neighborhood and root-sign conditions before this policy is accepted.
B1	Document or input form	Both consume one IR initially. Lift existing claim chains before designing a large grammar. Prefer a rendering if it matches reader benefits without authoring overhead.
B2	What is routine?	Instrument successful and failed discharge attempts under fixed contexts and budgets. A lexical tactic family is not a routine-discharge guarantee. Measure needed author hints separately.
B3	Analysis plumbing	Start with a composite locality/derivative theorem adapter. Compare it with smaller congruence interfaces on held-out local proofs; retain whichever exposes assumptions with less work.
B4	Roles and refactoring	Use stable project-owned role interfaces with typed wrappers, not raw upstream binder positions or names as permanent identities. Upgrade the upstream theorem and check the wrapper boundary.
B5	Notice dismissal	Statement-local or explicit project-convention dismissals, keyed to the relevant semantic expression and version. Mutating division or quantifier dependence must trigger re-review.

ID	Question area	Decision and discriminating artifact
B6	Replay unit	A step is the source-map unit; shared terms and certificates are storage units. Retained evidence is authoritative only for its pin. Compare script replay, certificate checking, and term replay costs.
B7	Provider outcomes	Keep diagnostics, fragment non-derivability, and checked negation separate. An abstraction list is not adequacy evidence. Adapt Leant and an algebra checker to distinct positive and negative transfer contracts.
B8	Value of structural synthesis	Measure compositions separately from single-lemma lookup. Keep the Haskell engine behind a Lean-owned adapter; do not rewrite it to justify the architecture.
B9	Published data holes	Publish the exact accepted witness and its specification. Keep rejected alternatives in an optional exploration log. Test whether alternatives improve understanding enough to expose them by default.
B10	Definition scope	Start with ordinary Lean definitions plus package templates and observation laws, not a general controlled-language compiler. The EGF package is the first counterexample-driven test.
B11	Contract ownership	Project wrappers own stable interfaces; profiles activate lexically. Conflicting operation conventions require explicit selection, while restrictive policies combine only if well typed. Test two independently authored packages.
B12	Stopping rule	A successful library or document view is acceptable. Continue a new input language only if the shared-library comparison shows an incremental authoring or reading benefit worth its maintenance.

B A discriminating conformance suite

This is a proposed integration suite, not a list of tests already executed against an implementation. The companion arithmetic program covers only the finite mathematical checks described in Appendix D. Each row needs a positive neighbor so that refusal of every query cannot pass the suite.

ID	Mutation	Required boundary and diagnostic
V1	Replace cast of natural subtraction by integer subtraction with no bound.	Missing non-truncation premise; no semantic rewrite.
V2	Cast natural $6 \div 4$ as rational $6/4$.	Exact-divisibility bridge fails.
V3	Reflect an equality modulo two into an integer equality.	No reflection contract; preserve-only route rejected.
V4	Add a competing representation route.	Same-anchor observations agree; selected presentation or method changes are visible.
V5	Use an index fact from another branch through a cache.	Origin/context mismatch; reject before use.
V6	Request coefficient N from a jet of length N .	Insufficient precision; request refinement.
V7	Differentiate a length- N jet and claim the same precision.	Derivative precision contract fails.
V8	Substitute a series with nonzero constant into the simple jet-composition rule.	Zero-constant premise or a stronger precision theorem is required.
V9	Interpret a formal identity analytically without convergence.	Missing analytic-evaluation bridge.
V10	Cancel $x - 1$ at $x = 1$ in a total real expression.	Support fails; exceptional branch remains.

ID	Mutation	Required boundary and diagnostic
V11	Cancel a nonzero zero divisor in a general ring.	Unit or regularity evidence required, not merely nonzero.
V12	Return a negative root for a nonnegative-root selection.	Full selection specification fails.
V13	Return one valid root as the complete root set.	Root-list completeness certificate missing.
V14	Change a Sturm remainder sign or use a root as an unhandled endpoint.	Exact checker rejects its recurrence or endpoint contract.
V15	Omit a singular parameter branch.	Coverage proof absent; result remains conditional.
V16	Promote a narrow interval to an exact decimal equality.	Wrong query specification; enclosure is not equality.
V17	Differentiate equality only at a point.	Neighborhood equality missing, before derivative search.
V18	Remove the last point of an inclusive reindexing.	Bijection coverage fails at the endpoint.
V19	Exchange limit and integral from pointwise convergence alone.	Exact analytic method hypotheses remain open.
V20	Accept an unused false checkpoint.	Document coverage fails even if the final theorem closes.
V21	Extract invalid source from a successful tactic's diagnostic.	Exact displayed-edit replay rejects the edit.
V22	Treat a reduction of two goals to one as theorem completion.	Progress receipt only; no completed root term.
V23	Finish after a changed instance or undo generation.	Stale origin rejected; no live-state mutation.
V24	Use fragment search exhaustion as a proof of negation.	No checked negation at the original target; report qualified failure.
V25	Reuse an accepted algorithm output without checked execution.	Algorithm theorem alone is insufficient for that output.
V26	Prove the target by an alias of the target in a benchmark.	Dependency exclusion invalidates the run.

C Minimal surface grammar and result schema

The grammar describes only the proposed stored document subset. The metavariable `leanExpr` denotes the host expression parser; `block` is indentation- delimited. There is no claim that this grammar implements all mathematical notation appearing in illustrative listings.

```

document    ::= item*
item        ::= object | claim | presentation | observation | selection
              | contextBlock | leanLeaf
object      ::= "object" ident ":" leanExpr "!=" leanExpr
claim       ::= "claim" ident binder* ":" leanExpr proof
proof       ::= "from" reference
              | "by" method arguments?
              | "by search" policy?
              | block
presentation ::= "present" reference "as" representation proof
observation  ::= "observe" ident "!=" leanExpr
              "compute using" method policy?
selection   ::= "select" ident ":" leanExpr
              "satisfying" leanExpr "using" method
contextBlock ::= "context" binder* block
leanLeaf    ::= "lean" verbatimLeanBlock
method      ::= qualifiedName
reference    ::= stableClaimOrObjectReference
policy      ::= "with" qualifiedProfileName

```

A checked computation record conceptually contains:

```

CheckedComputation
  origin          -- exact query, context, anchors, generation
  specification   -- expected output relation, not only a type
  output          -- exact accepted value or representation
  support         -- ordered instantiated guard telescope
  precision       -- observation restriction, when applicable
  coverage        -- only for a claimed complete case result
  evidence        -- term or checked certificate + soundness theorem
  execution_policy -- kernel/trace checking or explicit extra trust
  method_and_package_ids -- logical and operational versions
  source_map      -- assertions and observation nodes established

```

Neither an untrusted JSON field named `verified` nor a content hash can construct this record. The authoritative worker constructs it only after the required checks. Deserialization validates sizes, references, expression types, and schema versions before certificate execution.

D Companion checks and source provenance

D.1 Executed mathematical checks

The companion `reference_checks.py` uses Python’s exact integers and `Fraction`, without external packages. It checks the polynomial identity (1), validates the negative-remainder recurrence of the Sturm chain (2), and verifies its sign variations at the stated endpoints. It also checks finite instances of normalized EGF multiplication and derivative transport, Touchard recurrence, Spivey’s identity, and binomial-kernel orthogonality. Negative controls distinguish finite precision, truncated arithmetic, and cancellation at a zero denominator.

These checks independently catch errors in the displayed examples and executable arithmetic, not errors in a yet-unimplemented Lean checker. Their finite test ranges and results are recorded by the program in `reference-checks.json`. The delivered run passed 3,013 checks in ten groups. Universal claims in the article rest on the mathematical derivations, and an implementation must still formalize the corresponding contracts and checking algorithms.

D.2 Directly inspected files

Repository	Files and scope
Brainstorm	Both requested synthesis TeX files, including conclusions, mutual reviews, appendices, and next-round questions; pinned at 494be9d5.
ProveIt	<code>BinomialInversion.lean</code> , <code>AutonomousIteratedDeriv.lean</code> , and <code>BellShiftEGF.lean</code> ; the first 200 lines of <code>BellGeneratingFunctions.lean</code> ; pinned at 24ce8bd7.
Leant	<code>Verification.hs</code> , first 240 lines; <code>Fragment.hs</code> , first 160 lines; <code>Main.hs</code> , lines 4700–5030 containing suggestion generation and acceptance; pinned at 3a40904.

Repository links below are pinned to inspected revisions. Moving external documentation is not a ProveIt build pin. The source manifest records that distinction; the archive also contains rebuild instructions and exact-arithmetic checks.

References

- [1] Codex. *Mathematical Ideas, Checked Evidence: A unified evaluation of nine proof-language proposals*. Brainstorm, 5 September 2026. [Pinned TeX source](#). Especially the architecture, Leant review, view-composition discussion, and ten next-round questions.
- [2] Claude (Anthropic). *Nine Blueprints for a Mathematician’s Proof Language over Lean*. Brainstorm, 5 September 2026. [Pinned TeX source](#). Includes the corpus audit, annotated-theorem proposal, parallel-synthesis review, and twelve next-round questions.
- [3] Vladimir Reshetnikov and contributors. ProveIt, `Analysis/FabiusFunction/Lean/FabiusFunction/BinomialInversion.lean`. [Pinned source](#). Declarations `sum_lcc_neg_one_pow_choose_mul_choose`, `binomial_inversion`, and ring variants.
- [4] ProveIt, `Analysis/FabiusFunction/Lean/FabiusFunction/AutonomousIteratedDeriv.lean`. [Pinned source](#). Declaration `AutonomousODE.iteratedDeriv_eq_comp`.
- [5] ProveIt, `Analysis/FabiusFunction/Lean/FabiusFunction/BellShiftEGF.lean`. [Pinned source](#). Declarations `derivative_touchardExp`, `egfA_bell_add`, and `spivey`.
- [6] ProveIt, `Analysis/FabiusFunction/Lean/FabiusFunction/BellGeneratingFunctions.lean`. [Pinned source](#). First 200 lines, especially `egfA`, `coeff_egfA`, `egfA_mul`, and `derivative_egfA`.
- [7] Vladimir Reshetnikov and contributors. Leant, `src/Leant/Synth/Verification.hs`. [Pinned source](#). First 240 lines: callback receipts and quota-aware candidate verification.
- [8] Leant, `src/Leant/Synth/Fragment.hs`. [Pinned source](#). First 160 lines: fragment description, nominal applications, opaque atoms, and qualification of negative outcomes.
- [9] Leant, `src/Main.hs`. [Pinned source](#). Lines 4700–5030, particularly `suggestTactic`, `probeOnce`, `renderChain`, and the origin-sensitive cache.
- [10] Lean team. *Elaboration and Compilation*. Lean Language Reference. <https://lean-lang.org/doc/reference/latest/Elaboration-and-Compilation/>. Moving documentation consulted 5 September 2026 Pacific.
- [11] Lean team. *Validating a Lean Proof*. Lean Language Reference. <https://lean-lang.org/doc/reference/latest/ValidatingProofs/>. Moving documentation consulted 5 September 2026 Pacific; especially native computation, axiom inventories, independent rechecking, and statement comparison.
- [12] Mathlib contributors. *Mathlib.Tactic.Polyrith*. <https://github.com/leanprover-community/mathlib4/blob/master/Mathlib/Tactic/Polyrith.lean>. Current module documentation consulted 5 September 2026 Pacific. The discontinued external Sage path is not assumed as an available backend.
- [13] CoqEAL contributors. *The Coq Effective Algebra Library*. <https://github.com/rocq-community/coqeal>. Project documentation on abstract algebra, computation-oriented representations, and refinement; consulted 5 September 2026 Pacific.
- [14] Wenda Li, Grant Olney Passmore, and Lawrence C. Paulson. *Deciding Univariate Polynomial Problems Using Untrusted Certificates in Isabelle/HOL*. arXiv:1506.08238v2, 10 April 2018. <https://arxiv.org/html/1506.08238v2>. Primary paper describing certificate-guided univariate real-algebraic reasoning.
- [15] Wenda Li. *The Sturm–Tarski Theorem*. Archive of Formal Proofs, 2014. https://isa-afp.org/entries/Sturm_Tarski.html. Formalization entry and development history, consulted 5 September 2026 Pacific.